

Learning and Deducing Temporal Orders

Wenfei Fan^{1,2,3}, Resul Tugay², Yaoshu Wang¹, Min Xie¹, Muhammad Asif Ali¹

Shenzhen Institute of Computing Sciences, China¹, University of Edinburgh, UK² Beihang University, China³
wenfei@inf.ed.ac.uk, resul.tugay@ed.ac.uk, yaoshuw@sics.ac.cn, xiemin@sics.ac.cn, asif.ali@sics.ac.cn

ABSTRACT

This paper studies how to determine temporal orders on attribute values in a set of tuples that pertain to the same entity, in the absence of complete timestamps. We propose a creator-critic framework to learn and deduce temporal orders by combining deep learning and rule-based deduction, referred to as GATE (Get the LATEst). The creator of GATE trains a ranking model via deep learning, to learn temporal orders and rank attribute values based on correlations among the attributes. The critic then validates the temporal orders learned and deduces more ranked pairs by chasing the data with currency constraints; it also provides augmented training data as feedback for the creator to improve the ranking in the next round. The process proceeds until the temporal order obtained becomes stable. Using real-life and synthetic datasets, we show that GATE is able to determine temporal orders with F -measure above 83%, improving deep learning by 11.4% and rule-based methods by 47.4%.

1 INTRODUCTION

Real-life data keeps changing. As reported by Royal Mail, its customer data changes frequently: on average, “9590 households move, 1496 people marry, 810 people divorce, 2011 people retire and 1500 people die” every day in the UK [55]. It is estimated that inaccurate customer data costs organizations 6% of their annual revenues [55]. Outdated data incurs damage not only to Royal Mail. When the data at a search engine is out of date, a restaurant search may return a business that had closed three years ago. When the data about the condition of infrastructure assets is obsolete, it may delay the maintenance of network equipment and cause outage. Moreover, data-driven decisions based on outdated data can be worse than making decisions with no data [45]. Indeed, “as a healthcare, retail, or financial services business you cannot afford to make decisions based on yesterday’s data” [22]. Unfortunately, “58% of organizations make decisions based on outdated data” [22].

These highlight the practical need for determining the currency of data, *i.e.*, how up-to-date the information is. This is, however, highly nontrivial. Consider a set of tuples that pertain to the same entity, which are perhaps extracted from different data sources and are identified by entity resolution methods. **Their attribute values may become obsolete and inaccurate over the time.** Worse yet, only partial reliable timestamps are available. How can we determine the temporal orders on the attribute values, *i.e.*, for tuples t_1 and t_2 in the set, whether the value in the A -attribute of t_1 is more current than the value in the A -attribute of t_2 , denoted by $t_2 <_A t_1$?

Example 1: Consider customer records t_1 - t_6 shown in Figure 1, which have been identified to refer to the same person Mary. Each tuple t_i has attributes FN (first name), LN (last name), sex, address, marital status, job, kids and SZ (shoe size). **Some attribute values of these tuples have become stale** since Mary’s data changes over the years, *e.g.*, her job, address and last name have changed 4 times, five times and twice, respectively. Only some attribute values might be

	FN	LN	sex	address	status	job	kids	SZ
t_1, e_1 :	Mary	Goldsmith	F	19 xin st	single	n/a	-	5
t_2, e_1 :	Mary	Taylor	F	6 gold plaza	married	journalist	1	6
t_3, e_1 :	Mary	Taylor	F	19 mall st	espoused	assoc editor	2	7
t_4, e_1 :	Mary	Taylor	F	7 ave	divorced	chief editor	2	7
t_5, e_1 :	Mary	Taylor	F	7 ave	detached	chief editor	2	7
t_6, e_1 :	Mary	Goldsmith	F	6 const. ave	wed	producer	3	7

Figure 1: Customer records dataset

associated with reliable timestamps, *e.g.*, the timestamp of t_5 [job] and t_6 [job] are 2016 and 2019 (not shown), respectively, indicating that at that time, the values are up-to-date. In the absence of complete timestamps, it is hard to know whether $t_2 <_{LN} t_6$, *i.e.*, whether the value of t_2 [LN] is more up-to-date than t_6 [LN]? Moreover, what Mary’s current job title is, *e.g.*, whether $t_i <_{job} t_6$ for $i \in [1, 5]$? □

One may want to approach this problem by training a ranking model that sorts objects according to their degrees of relevance, preference or importance [47]. The state-of-the-art systems in this regard employ deep learning [7, 33, 52, 58] or reinforcement learning [34, 63], and have been used in search engine [9] and machine translation [21, 65]. By means of ranking models one can learn temporal orders and decide whether $t_1 <_A t_2$ for all tuples t_1, t_2 and attribute A . However, it is hard to justify whether the ranking conforms to the temporal orders on the values in the real world. For data-driven decision making, we need to ensure that the learned orders are reliable. Moreover, these approaches are not capable of explaining the ranking of objects that follow a complex interlinked ordering structure (*e.g.*, [address](#) in Figure 1).

Another approach is to employ logic rules, *e.g.*, currency constraints [18, 20, 25, 27, 40, 44, 61]. These rules may specify the monotonicity and comonotonicity of data values, and help us deduce temporal orders. Taking the tuples from Example 1 as an example, the shoe sizes of the same person typically monotonically increases (before 20 years old), and **the address of a person may be** associated with the marital status (marital status changes from single to married, not the other way around). As will be seen in Section 2, using currency constraints one can deduce that $t_2 <_{SZ} t_3$ and $t_1 <_{address} t_2$. However, it is hard to find enough rules to deduce relative orders on each and every pair of attribute values. For instance, we cannot find rules to decide whether $t_2 <_{LN} t_6$. Besides, it is also hard to generalize logic rules to handle lexically different but semantically similar attribute values, since tuples might be extracted from different source (*e.g.*, attribute status in Figure 1: married vs. wed, espoused and divorced vs. detached).

Neither deep/machine learning models nor logic rules work well when they are used separately. A natural question is whether it is possible to combine ML models and logic rules in a uniform framework, such that we can learn temporal orders for attribute values, and use the rules to deduce more orders and justify the ranking? How well can the uniform framework improve the accuracy of the learning models and rules when being used alone?

Contributions & organization. This paper makes a first effort to

answer these questions, all in the affirmative.

(1) *Temporal orders* (Section 2). We define the notion of temporal orders $t_1 <_A t_2$ and $t_1 \leq_A t_2$ on attributes, and formulate the problem for determining temporal orders. We also review the currency constraints (CCs) of [27], for deducing temporal orders by our framework. We show that CCs can specify interesting temporal properties, e.g., the monotonicity, comonotonicity and transitivity.

(2) *GATE* (Section 3). We propose GATE (Get the *l*AtESt), a creator-critic framework to determine temporal orders by combining deep learning and discovered CCs. GATE iteratively invokes a creator to rank the temporal orders on attribute values, followed by the critic that validates the ranking of the creator and deduces more ranked pairs via logic deduction. The critic also produces augmented training data for the creator to improve its ranking in the next round. This process proceeds for the creator and critic to mutually enhance each other, until the temporal order obtained cannot be improved.

(3) *Creator* (Section 4). We propose a deep learning model underlying the creator of GATE, to learn temporal orders on attribute values. It departs from previous ranking models in that it learns the temporal orders based on contexts (attribute correlations), calculates the confidence of the ranking, and employs an attribute-centric adaptive pairwise ranking strategy. It also projects attribute embedding to a low-dimensional space that preserves chronological orders.

(4) *Critic* (Section 5). The critic complements GATE with discovered rules (CCs). We show how it justifies the ranked pairs learned by the creator, and (incrementally) deduces latent temporal orders with the chase [57] using CCs. The chase has the Church-Rosser property (cf. [2]), i.e., it guarantees to converge at the same result no matter in what orders the CCs are applied. Moreover, the critic augments the training data for the creator to improve its model in the next round.

(5) *Experimental study* (Section 6). Using real-life and synthetic data, we find: (a) The *F*-measure of GATE on dataset Career [37] is 0.866, versus 0.35 (resp. 0.36) by rule-based UncertainRule [40] (resp. Improve3C [17]), and 0.54 (resp. 0.53) by ML-based RANK_{Bert} [51] (resp. Ditto_{Rank}). (b) On average, GATE is 47.4% and 11.3% more accurate than the critic and the creator, respectively, verifying the effectiveness of combining deep learning and logic rules. (c) GATE is feasible in practice; it only takes 7 rounds to terminate on a real-life dataset of 1,983,698 tuples, with a single machine.

In summary, we study a new problem for determining temporal orders on attribute values. We propose the first framework that combines deep learning and logic deduction to enhance each other. We also provide a novel ranking model to learn temporal orders, and extend the chase to justify and improve the learner.

Related work. We categorize the related work as follows.

Rule-based methods. Currency constraints (CCs) were first studied in [26, 27] by employing partial currency orders. The work was extended in [24, 25] for conflict resolution, by considering both CCs and conditional functional dependencies (CFDs) [23]. A class of currency repairing rules (CRRs) was proposed in [39], which combines logic rules and statistics. A two-step approach was developed in [16] to determining the currency of dynamic data, by means of a topological graph to represent currency orders. A class of uncertain currency rules was supported by UncertainRule [40] that

considers both temporal orders and data certainty. Improve3C [17] and Imp3C [18] are 4-step data cleaning frameworks, including data consistency, completeness and currency, which use a metric to repair noisy data using CFDs [23] and currency orders in particular. There has also been work [20] on parallel incremental updating algorithms by employing traditional currency rules.

This work differs from the prior work in the following. (a) To the best of our knowledge, we make the first effort to combine deep learning and logic rules for determining currency, while the prior work at most extends rules with statistical data. (b) We propose a deep-learning model for inferring temporal orders and a logic deduction model for incrementally deriving more ranked pairs. (c) We develop a creator-critic framework such that deep-learning models and rule-based methods can enhance each other.

ML models. There have also been efforts on inferring temporal orders of events extracted from text by ML models [6, 8, 32, 49, 50, 59]. A related topic is to incorporate temporal information into knowledge graphs, e.g., for temporal link predictions [19, 56] and temporal reasoning [15, 60, 64]. This line of work often focuses on how to embed temporal information into ML models, for inference that varies over time. Moreover, learning temporal orders has been modeled as a learning-to-rank problem (including pointwise, pairwise and listwise ranking), where global orders are learned for currency attributes. Temporal problems have also been studied in, e.g., search and recommendation [28, 62], and information retrieval [46].

Different from the prior work, we propose a creator-critic framework that not only learns temporal orders via deep learning, but also adopts logic rules to deduce new ones and help ML models learn better. Our creator learns temporal orders based on contexts, calculates the confidence of the ranking, and incrementally improves its model with augmented training data from the critic. The previous ML models can also be embedded into our framework.

Entity resolution (ER). ER has also been studied by incorporating temporal information. Temporal clustering methods were proposed in [42] based on a concept of time decay, to improve record linkage with timestamps. A method was developed in [10] to dynamically assess and adjust similarities among records based on their attributes and temporal differences. A transition model was developed in [38] to estimate the timeliness of values (i.e., the probability that a value changes to other values) and show how to link tuples in a correct time period. To cope with incorrect timestamps, [61] adopts matching dependencies (MDs) and data currency constraints to capture the attributes that were changed over time in entity resolution.

This work focuses on deducing temporal orders of each entity by combining deep learning models and logical rules, while the prior work mainly adopts temporal information to solve the ER problem.

2 DETERMINING TEMPORAL ORDERS

In this section we first formulate temporal orders and the problem for determining temporal orders (Section 2.1). We then review currency constraints (CCs) of [27] and show that such rules are able to express monotonicity, comonotonicity and transitivity (Section 2.2).

2.1 The Problem

Consider a relation schema $R = (EID, A_1, \dots, A_n)$, where A_i is an attribute ($i \in [1, n]$), and EID is an entity id as introduced by

Codd [13]. For a tuple t of schema R , we denote by $t[A]$ the value in the A -attribute of t , and by $t[\text{EID}]$ the entity represented by tuple t . We assume that each tuple is associated with a tuple id, denoted by tid . We refer to a relation D of R as a *normal instance* D of R .

Temporal orders and temporal instances. An *entity instance* is (D_e, T_e) , where (a) D_e is a normal instance of schema R such that all tuples t_1 and t_2 in D_e refer to the same real-life entity e and hence, $t_1[\text{EID}] = t_2[\text{EID}]$; and (b) T_e is a partial function that associates a timestamp $T_e(t[A])$ with the A -attribute of a tuple t in D_e . We refer to (D_e, T_e) as an *entity instance pertaining to e* .

Here the timestamp indicates that at the time $T_e(t[A])$, the A -attribute value of tuple t is correct and up-to-date; it does not necessarily refer to the time when $t[A]$ was created or last updated. If $T_e(t[A])$ is undefined, a reliable timestamp is not available for $t[A]$. **Note that we do not assume a timestamp for the entire tuple, since we often find only parts of a tuple to be correct and up-to-date.**

Intuitively, an entity instance extends a normal instance with available timestamps. Its tuples may be extracted from a variety of data sources, and are identified to refer to the same entity e via entity resolution. In the same tuple t , $t[A]$ and $t[B]$ may bear different timestamps (or even no timestamp) for different A and B .

Temporal orders. A *temporal order* on attribute A of D_e is a partial order \leq_A such that for all tuples t_1 and $t_2 \in D_e$, $t_2 \leq_A t_1$ if the value in $t_1[A]$ is at least as current as $t_2[A]$. We also use a strict partial order $t_2 <_A t_1$ if $t_1[A]$ is more current than $t_2[A]$. **To simplify the discussion we focus on \leq_A in the sequel; $<_A$ is handled analogously.**

In particular, if $T_e(t_1[A])$ and $T_e(t_2[A])$ are both defined and if $T_e(t_2[A]) \leq T_e(t_1[A])$, i.e., when timestamp $T_e(t_1[A])$ is no earlier than $T_e(t_2[A])$, then $t_2 \leq_A t_1$, i.e., $t_1[A]$ is confirmed at a later timestamp and is thus considered at least as current as $t_2[A]$.

Temporal order \leq_A is represented as a set of tuple pairs such that $(t_2[\text{tid}], t_1[\text{tid}]) \in \leq_A$ iff $t_2 \leq_A t_1$. We write $(t_2[\text{tid}], t_1[\text{tid}])$ as (t_2, t_1) if it is clear in the context. Note that the same value may bear different timeliness in different tuples, e.g., *Mary's marital status changed from married (t_2) to divorced (t_4) to married (t_7 , not shown in Figure 1). While $t_2[\text{status}] = t_7[\text{status}]$, $t_2 <_{\text{status}} t_7$. Here $t_2 <_{\text{status}} t_7$ ranks the timeliness of the status-attributes of tuples t_2 and t_7 , not values (married vs. married) detached from the tuples.*

We say that a temporal order \leq_A^1 *extends* \leq_A^2 , written as $\leq_A^2 \subseteq \leq_A^1$, if for all tuples t_1, t_2 in D_e , if $t_2 \leq_A^2 t_1$ is defined, then so is $t_2 \leq_A^1 t_1$. That is, \leq_A^1 includes all tuple pairs in \leq_A^2 and possibly more.

Temporal instances. A *temporal instance* D_t of R is given as $(D, \leq_{A_1}, \dots, \leq_{A_n}, T)$, where each \leq_{A_i} is a temporal order on A_i ($i \in [1, n]$), $D = \bigcup_{i \in [k]} D_{e_i}$, $T = \bigcup_{i \in [k]} T_{e_i}$, and for all $i \in [1, k]$, (D_{e_i}, T_{e_i}) is an entity instance of R . Here tuples t_1 and t_2 in D are *compatible* under \leq_A if they pertain to the same entity, i.e., $t_1[\text{EID}] = t_2[\text{EID}]$.

Intuitively, D_t is a collection of entity instances, such that each (D_{e_i}, T_{e_i}) pertains to the same entity e_i . We do not rank the currency of tuples if they refer to different entities. A temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ is said to *extend* another temporal instance $D'_t = (D, \leq'_{A_1}, \dots, \leq'_{A_n}, T)$ if for all $i \in [1, n]$, \leq_{A_i} extends \leq'_{A_i} .

Problem statement. We study the problem for *determining the temporal orders* of a temporal instance, stated as follows.

- **Input:** A temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$.

- **Output:** An extended temporal instance $D'_t = (D, \leq'_{A_1}, \dots, \leq'_{A_n}, T)$ such that for all $i \in [1, n]$, (a) \leq'_{A_i} extends \leq_{A_i} and (b) \leq'_{A_i} is a total order on all compatible t_1 and t_2 with $t_1[\text{EID}] = t_2[\text{EID}]$.

Intuitively, **our goal** is to extend \leq_{A_i} such that for all tuples t_1 and t_2 in D , if t_1 and t_2 refer to the same entity e , then we can decide which of $t_1[A_i]$ and $t_2[A_i]$ is more up-to-date. As an immediate consequence, we can deduce the latest value for all attributes. **Note that we define a total order on each attribute, not a global order on tuples. In other words, the total orders on different attributes can be different, and the latest values for different attributes of entity e may come from multiple tuples in D_e . This said, the deduction of temporal orders on one attribute may help the deduction on other attributes, via correlation expressed as currency constraints (see below).**

2.2 Currency Constraints

We next review the class of currency constraints proposed in [27]. The predicates over a relation schema R are defined as:

$$p ::= t[A] \oplus c \mid t_1[A] \oplus t_2[A] \mid t_1 \leq_A t_2,$$

where t, t_1, t_2 are tuple variables denoting tuples of R , A is an attribute of R , c is a constant, \oplus is an operator from $\{=, \neq, >, \geq, <, \leq\}$; $t[A] \oplus c$ and $t_1[A] \oplus t_2[A]$ are defined on attribute values, while $t_1 \leq_A t_2$ compares the timeliness of $t_1[A]$ and $t_2[A]$. In particular, $t_1[\text{EID}] = t_2[\text{EID}]$ says that t_1 and t_2 refer to the same entity.

A *currency constraint* (CC) over schema R is defined as follows:

$$\varphi = X \rightarrow p_0,$$

where X is a conjunction of predicates over R with tuple variables t_1, \dots, t_m , and p_0 has the form $t_u \leq_{A_i} t_v$ for $u, v \in [1, m]$. We refer to X as the *precondition* of φ , and p_0 as the *consequence* of φ .

As defined in [27], a CC can be equivalently expressed as a universal first-order logic sentence of the following form:

$$\varphi = \forall t_1, \dots, t_m \left(\bigwedge_{j \in [1, m]} (t_1[\text{EID}] = t_j[\text{EID}]) \wedge X \rightarrow t_u \leq_{A_i} t_v \right).$$

Semantics. Currency constraints are defined over temporal instances $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ of R . A *valuation* of tuple variables of a CC φ in D_t , or simply a *valuation* of φ , is a mapping h that instantiates variables t_1, \dots, t_m with tuples in D that refer to the same real-world entity, as required by $t_1[\text{EID}] = t_j[\text{EID}]$.

A valuation h *satisfies* a predicate p over R , written as $h \models p$, if the following is satisfied: (1) if p is $t[A] \oplus c$ or $t_1[A] \oplus t_2[A]$, then it is interpreted as in tuple relational calculus following the standard semantics of first-order logic [2]; and (2) if p is $t_1 \leq_A t_2$, then $t_2[A]$ is at least as current as $t_1[A]$ i.e., (t_1, t_2) is in \leq_A .

For a conjunction X of predicates, we write $h \models X$ if $h \models p$ for all p in X . A temporal instance D_t *satisfies* CC φ , denoted by $D_t \models \varphi$, if for all valuations h of $X \rightarrow p_0$ in D_t , if $h \models X$ then $h \models p_0$.

Properties. Currency constraints are able to specify interesting temporal properties. Below we exemplify some properties.

Monotonicity. A temporal order \leq_A over relation schema R is *monotonic* if for any tuples t_1 and t_2 of R that refer to the same entity, if $t_1[A] \leq t_2[A]$ then $t_1 \leq_A t_2$. For instance, consider the SZ (shoe size) attribute of the customer relation of Example 1. Then \leq_{SZ} is monotonic. It can be expressed as the following CC:

$$\varphi_1 = t_1[\text{SZ}] \leq t_2[\text{SZ}] \rightarrow t_1 \leq_{\text{SZ}} t_2.$$

As another example, marital status only changes from single to

Symbols	Notations
$R = (A_1, \dots, A_n)$	relation schema
(D_e, T_e)	instance of schema R pertaining to entity e
temporal order $t_1 \leq_A t_2$	t_2 is at least as current as t_1 in attribute A
$\text{conf}(t_1 \leq_A t_2)$	the confidence of $t_1 \leq_A t_2$
$(D, \leq_{A_1}, \dots, \leq_{A_n}, T)$	temporal instance of R
$\varphi = X \rightarrow p_0$	currency constraint
h	valuation of φ in a temporal instance
Γ	the ground truth
D_{aug}	the augmented training data
GATE	Get the LATEST

Table 1: Notations

married, not the other way around. This is expressed as CC:
 $\varphi_2 = t_1[\text{status}] = \text{"single"} \wedge t_2[\text{status}] = \text{"married"} \rightarrow t_1 \leq_{\text{status}} t_2$.

With slight abuse of terminologies by considering timestamps as an “attribute” associated with attribute A , we have:

$$\varphi_3 = T_e(t_1[A]) \leq T_e(t_2[A]) \rightarrow t_1 \leq_A t_2,$$

since $t_2[A]$ is confirmed at a later timestamp.

Comonotonicity. For attributes A and B of R , we say \leq_A and \leq_B are *comonotonic* in a temporal instance D_t of R if for all tuples t_1 and t_2 in D_t that refer to the same entity, $t_1 \leq_A t_2$ if and only if $t_1 \leq_B t_2$.

Intuitively, \leq_A and \leq_B are comonotonic if the two are correlated such that when one is updated, the other will also change. For instance, \leq_{status} and \leq_{address} are often comonotonic: when the marital status of a person changes from single to married, *he/she may move to a larger house*. This can be expressed as a CC:

$$\varphi_4 = t_1 \leq_{\text{status}} t_2 \rightarrow t_1 \leq_{\text{address}} t_2.$$

Transitivity. Transitivity can also be expressed for any attribute A :

$$\varphi_5 = t_1 \leq_A t_2 \wedge t_2 \leq_A t_3 \rightarrow t_1 \leq_A t_3.$$

Correlating different attributes. One can correlate multiple different attributes to capture implicit temporal ordering. For example, marital status may change from “married” to “divorced” and further from “divorced” to “married” again. To deduce an order on $t[\text{status}]$, we can use the number of kids as an additional condition:

$$\varphi_6 = t_1[\text{status}] = \text{"married"} \wedge t_2[\text{status}] = \text{"divorced"} \wedge t_1[\text{kids}] < t_2[\text{kids}] \rightarrow t_1 \leq_{\text{status}} t_2.$$

Deduction. Making use of CCs, we can deduce certain temporal orders, e.g., from φ_1 - φ_6 and Figure 1, we can deduce the following:

- $t_1 \leq_{\text{status}} t_2$ by φ_2 , $t_2 \leq_{\text{status}} t_4$ by φ_6 and $t_1 \leq_{\text{status}} t_4$ by φ_5 ;
- $t_1 \leq_{\text{address}} t_2$ by φ_4 ; hence $t_2[\text{address}]$ is more current for Mary;
- $t_1 \leq_{\text{SZ}} t_2 \leq_{\text{SZ}} t_3$ by φ_1 ; hence Mary’s current shoe size is 7.

However, we cannot determine whether $t_2 \leq_{\text{LN}} t_6$ or $t_6 \leq_{\text{LN}} t_2$ by deduction with the currency constraints φ_1 - φ_6 .

Discovery of currency constraints. As noted in [27], CCs can be considered as a special case of denial constraints (DCs) introduced in [4], extended with temporal orders \leq_A . Instead of being manually constructed, several methods have been developed for discovering DCs automatically, e.g., FastDC [11], Hydra [5], DCFinder [53] and ADCMiner [48]. We extend them for CCs discovery in [1].

Remark. We remark that the discovery algorithm is executed once for each relation schema R by sampling its temporal instances. The set Σ of discovered currency constraints is then applied to different temporal instances. In other words, GATE does not have to discover currency constraints for each input temporal instance.

The notations of the paper are summarized in Table 1.

3 GATE: A CREATOR-CRITIC SYSTEM

In this section, we propose a creator-critic framework for determining temporal orders, and develop system GATE to implement it. A unique feature of GATE is its *combined* use of deep learning and logic deduction. Below we start with the architecture of GATE, and then present its overall workflow, with termination guarantee.

Architecture. The ultimate goal of GATE is to obtain a total order \leq_A for each attribute A . As shown in Figure 2, GATE first discovers a set Σ of CCs on D_t *offline* for performing logic deduction. Then it takes a temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ as input, and learns and deduces more temporally ranked pairs for D_t *online*. For simplicity, we assume w.l.o.g. that D consists of a single entity instance D_e , i.e., all tuples in D pertain to the same entity e and thus their attributes can be pairwise compared; the methods of this paper can be readily extended to D_t with multiple entity instances.

More specifically, the learning and deducing process in GATE *iteratively* executes two phases, namely, creator and critic, as follows.

(1) **Creator.** In this phase, GATE (incrementally) trains a ranking model $\mathcal{M}_{\text{rank}}$ via deep learning. By taking D_t and the augmented training data D_{aug} (see its definition shortly) from the critic as input, it predicts new orders for extending each \leq_A of D_t . Our model has three features: (a) For each tuple t in D , the embedding for $t[A]$ is created using both $t[A]$ and other correlated values in t , so that $t[A]$ can be ranked comprehensively. (b) The attribute embeddings [14] are arranged to preserve chronological orders. (c) We adopt an *attribute-centric adaptive pairwise ranking* strategy in $\mathcal{M}_{\text{rank}}$, so that the ranking result can be justified semantically. Moreover, the temporal instance D_t will also be extended based on D_{aug} .

Given an attribute A , we associate each $(t_1, t_2) \in \leq_A$ with a *confidence*, denoted by $\text{conf}(t_1 \leq_A t_2)$, indicating how likely $t_1 \leq_A t_2$ holds. If $t_2[A]$ has a later timestamp than $t_1[A]$, then $\text{conf}(t_1 \leq_A t_2)$ is 1. If $t_1 \leq_A t_2$ is predicted by $\mathcal{M}_{\text{rank}}$, its confidence is from 0 to 1. We only consider (t_1, t_2) predicted by $\mathcal{M}_{\text{rank}}$ with $\text{conf}(t_1 \leq_A t_2) \geq \delta$, where δ is a predefined threshold, as *candidate pairs* to be extended to \leq_A . We denote the set of all candidate pairs of \leq_A by \leq_A^M .

The input and output of the creator are as follows:

- **Input:** A temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ and the augmented training data D_{aug} from the critic.
- **Output:** An extended temporal instance $D'_t = (D, \leq'_{A_1}, \dots, \leq'_{A_n}, T)$ based on D_{aug} and the predicted orders $(\leq'_{A_1}, \dots, \leq'_{A_n})$.

(2) **Critic.** In this phase, the critic of GATE justifies and deduces more temporally ranked pairs, by applying CCs in Σ via the *chase* [57]. Denote the result of chasing by \leq_A^Σ . Depending on the validity of \leq_A^Σ , we construct an augmented training data, denoted by D_{aug} , containing the ranked pairs justified (resp. conflicts caught by CCs); D_{aug} will be fed back to the creator, for the next round of model learning, so that the creator can learn from more unseen data and get higher accuracy *iteratively*. Specifically, D_{aug} is a temporal instance extended with a validity flag f_{valid} , i.e., $D_{\text{aug}} = (D, \leq'_{A_1}, \dots, \leq'_{A_n}, T, f_{\text{valid}})$: (a) if f_{valid} is true, \leq_A^Σ is valid and all temporal order deduced by the chase will be added to D_{aug} ; and (b) otherwise, the chasing is invalid, i.e., both $(t_1, t_2) \in \leq_A$ and $(t_2, t_1) \in \leq_A$ are deduced, and either $t_1 <_A t_2$ and $t_2 <_A t_1$. In this case, these two




Figure 2: Architecture of GATE

conflicting orders will be added to D_{aug} , and the creator will be asked to resolve this conflict, by revising its model accordingly.

Formally, the input and output of the critic are as follows:

- *Input:* A temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$, the predicted temporal orders $(\leq_{A_1}^M, \dots, \leq_{A_n}^M)$ and a set Σ of CCs.
- *Output:* Augmented data $D_{\text{aug}} = (D, \leq_{A_1}', \dots, \leq_{A_n}', T, f_{\text{valid}})$.

The novelty of the critic consists of (a) the deduction using the chase, and (b) an efficient algorithm implementing the chase.

Workflow. As shown in Figure 3, GATE takes a temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ and a set Σ of CCs discovered offline as input, and it outputs an extended temporal instance $D_t' = (D, \leq_{A_1}', \dots, \leq_{A_n}', T)$ with a total order \leq_A' defined for each attribute A .

GATE first initializes the augmented training data $D_{\text{aug}} = (D, \leq_{A_1}', \dots, \leq_{A_n}', T, f_{\text{valid}} = \text{true})$ (Line 1, details omitted) for the first round of GATE, by deducing its initial temporal orders \leq_A' via CCs in Σ whose preconditions do not involve timeliness comparison, e.g., we can create temporal orders for those tuples with available timestamps using φ_3 . The initialization can be done efficiently by [30].

Then GATE iteratively executes the creator and the critic in rounds. In the i -th round, (a) the creator *incrementally* trains $\mathcal{M}_{\text{rank}}$ (Line 4, see Section 4) based on the augmented training data D_{aug} returned by the critic in the $(i-1)$ -th round. Besides, the creator extends the temporal instance D_t (Line 5, see Section 4) based on D_{aug} , by possibly revising its model to resolve the conflicts. The creator then predicts new temporal orders $(\leq_{A_1}^M, \dots, \leq_{A_n}^M)$ (Line 6) based on the fine-tuned model, which are candidate pairs (t_1, t_2) with confidences at least δ ; and (b) the critic deduces more ranked pairs $(\leq_{A_1}^\Sigma, \dots, \leq_{A_n}^\Sigma)$ by chasing with CCs in Σ (Line 8). According to the result of chasing, the critic constructs augmented training data D_{aug} via procedure ConstructAugmented (Line 9, see Section 5); this D_{aug} is fed back to the creator for the next round processing.

Finally, when D_t no longer changes (Line 10-11), the iteration ends. If D_t is still not a temporal instance with total orders defined on all attributes, we extend it using procedure Extend (Line 12), such that for each pair (t_1, t_2) , one of (t_1, t_2) and (t_2, t_1) learned with a higher confidence is in \leq_A , until each \leq_A becomes a total order.

Termination. We prove that eventually, GATE will terminate (Line 10), i.e., with more iterations, D_t is gradually extended for more temporal orders and finally, becomes stable and does not change.

Theorem 1: GATE is guaranteed to terminate. \square

Proof. We prove Theorem 1 by showing that each temporal order \leq_A will be stable after certain rounds, since (a) once a temporal order $t_1 \leq_A t_2$ is stable, i.e., if (t_1, t_2) is added to \leq_A of D_t via procedure Extend (Line 5), it will not be removed from \leq_A in the subsequent rounds; and (b) the number of stable pairs in \leq_A is

Input: A temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$, and a set Σ of CCs.

Output: An extended temporal instance $D_t' = (D, \leq_{A_1}', \dots, \leq_{A_n}', T)$ such that for all $i \in [1, n]$, (a) \leq_{A_i}' extends \leq_{A_i} and (b) \leq_{A_i}' is a total order.

1. $D_{\text{aug}} := \text{Initialize}(D_t, f_{\text{valid}} = \text{true}, \Sigma)$; Initialize the ranking model $\mathcal{M}_{\text{rank}}$;
2. **while** true **do**
3. /* The Creator of GATE */
4. Train $\mathcal{M}_{\text{rank}}$ incrementally based on D_{aug} ;
5. $D_t := \text{Extend}(D_t, D_{\text{aug}})$;
6. $(\leq_{A_1}^M, \dots, \leq_{A_n}^M) :=$ the predicted temporal orders by $\mathcal{M}_{\text{rank}}$;
7. /* The Critic of GATE */
8. $(\leq_{A_1}^\Sigma, \dots, \leq_{A_n}^\Sigma) := \text{Chase}(D_t, \Sigma)$;
9. $D_{\text{aug}} := \text{ConstructAugmented}(D, \leq_{A_1}^M, \dots, \leq_{A_n}^M, \leq_{A_1}^\Sigma, \dots, \leq_{A_n}^\Sigma, T)$;
10. **if** D_t no longer changes **then**
11. **break**;
12. $D_t = \text{Extend}(D_t, \mathcal{M}_{\text{rank}})$;
13. **return** D_t ;

Figure 3: Workflow of GATE

strictly increasing, which is upper-bounded by $\mathcal{O}(|D|(|D| - 1))$. \square

Example 2: Continuing with Example 1, assume that $\Sigma = \{\varphi_1 - \varphi_6\}$ and D_t has empty temporal orders. We first initialize the training data D_{aug} by applying CCs in Σ that do not compare timeliness in their preconditions, e.g., we can deduce $t_5 \leq_{\text{job}} t_6$ by φ_3 . Suppose that after training $\mathcal{M}_{\text{rank}}$ based on the initial D_{aug} , no tuple pair predicted by $\mathcal{M}_{\text{rank}}$ is at least δ -confident in the first round. The creator outputs empty \leq_A^M for each A due to the lack of information. Then by applying the CCs in Σ , the critic can deduce new temporal orders \leq_A^Σ , e.g., $t_1 \leq_{\text{address}} t_2$ by φ_4 and $t_1 \leq_{\text{status}} t_4$ by φ_5 . Both \leq_A^M and \leq_A^Σ will be used to construct the augmented training data D_{aug} , based on which the creator extends D_t and incrementally trains $\mathcal{M}_{\text{rank}}$ in the second round. This time the creator might be able to predict confident temporal orders, e.g., $t_1 \leq_{\text{LN}} t_2$, with the augmented training data D_{aug} . The iteration continues until D_t does not change anymore. Each temporal order \leq_A in D_t will be extended to a total order by $\mathcal{M}_{\text{rank}}$ if it is still not total. \square

Remark. We adopt a confidence threshold δ to ensure the reliability of ML predictions, so that only reliable temporal orders are considered. However, if confident orders cannot be decided for the lack of initial information, we could optionally invite user inspection to ensure the correctness of a few initial ranked pairs, from which more reliable orders can be iteratively deduced/learned.

We present the creator and critic in Sections 4 and 5, respectively.

4 CREATOR

In this section, we develop the creator underlying GATE. Given a temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ and the augmented training data D_{aug} (from the critic), the creator (a) trains a ranking model $\mathcal{M}_{\text{rank}}$ to predict new temporal orders and (b) extends each \leq_A of D_t based on D_{aug} . Below we first review popular ranking models. We then present our ranking model $\mathcal{M}_{\text{rank}}$, how it is (incrementally) trained based on D_{aug} and how it extends D_t .

Review on ranking models. *Learning to Rank* [46] aims to learn a ranking model so that objects can be ranked based on their degrees of relevance, preference, or importance (in our setting, timeliness).

Existing learning-to-rank methods can be classified into three types: pointwise, pairwise and listwise (surveyed in [46]). We adopt the pairwise ranking setting, since (a) pairwise ranking is consistent with temporal orders, which is a set of tuple pairs on which partial

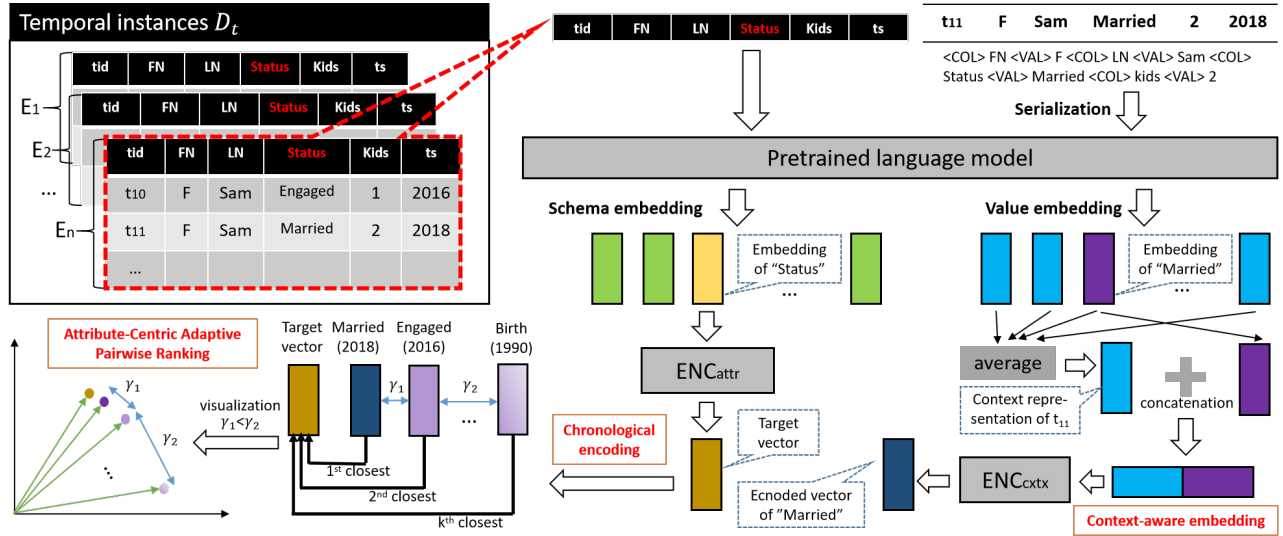


Figure 4: The Network Architecture of Creator

orders are defined; and (b) by the transitivity of temporal orders, pairwise ranking can eventually learn a total order for all attributes.

Challenges. One naturally wants to adopt an existing model. However, it is hard to directly apply one, since the unique features of temporal orders are not well considered, resulting in poor performance.

(1) *Attribute correlation.* Due to the correlated nature of temporal order, we often need to reference other attributes to determine the orders of a given attribute. Moreover, since a value may change back and forth (e.g., the marital status changes from “married” to “divorced”, and back from “divorced” to “married”), it is hard to determine the up-to-date value based on a single attribute only. This motivates us to explicitly consider attribute correlation.

(2) *Limitation of embedding models.* To determine timeliness, care should be taken for lexically different but semantically similar values (e.g., “baby” vs. “birth” and “dead” vs. “expired” for status). Although existing embedding models (e.g., ELMo [54] or Bert [14]) are widely adopted, they cannot be directly used here since they are not trained to organize data chronologically. This calls for a novel encoding mechanism to project the embeddings to a dense space, which preserves timeliness and meanwhile, maintains data semantics.

(3) *Adaptive margin.* Existing ranking strategies do not consider real-life characteristics of timeliness, e.g., the timespan for a person’s status to move from “birth” to “engaged” is typically longer than from “engaged” to “married” (Figure 4). Instead of ranking status with a *fixed* margin as most existing strategies did, we need a new methodology to embed values using *adaptive* margins, to conform to their real-life behaviors and justify the semantic of ranking.

Model overview. We propose a ranking model to tackle the above challenges, whose novelty includes (a) a context-aware scheme that embeds each value along with other correlated values, (b) an encoding mechanism to re-organize the embeddings in a chronological manner, and (c) an attribute-centric adaptive ranking strategy.

As shown in Figure 4, our ranking model $\mathcal{M}_{\text{rank}}$ takes the current D_t as input, and outputs new ranked pairs, where each (t_1, t_2) is associated with a *confidence*, indicating how likely $t_1 \leq_A t_2$ holds.

Our ranking model consists of three stages as follows: context-aware embedding, chronological encoding and order prediction.

(1) $\mathcal{M}_{\text{rank}}$ first builds a context-aware embedding for each attribute value using pre-trained language models (ELMo [54] or Bert [14]), in a way that information of correlated values is also embedded.

(2) Based on the embeddings, $\mathcal{M}_{\text{rank}}$ encodes a target vector ϕ_A for attribute A , via *non-linear transformation* (see below). Similarly, a value vector $\phi_{t[A]}$ is encoded for each A -attribute value of tuple t , so that (a) if $t[A]$ is more current, $\phi_{t[A]}$ is closer to ϕ_A and (b) the gap between $\phi_{t[A]}$ and ϕ_A is trained adaptively, to reflect real semantics.

(3) Finally, given $\phi_{t_1[A]}$ and $\phi_{t_2[A]}$ of t_1 and t_2 , $\mathcal{M}_{\text{rank}}$ predicts the order, i.e., whether $t_1 \leq_A t_2$ holds with high enough confidence.

Design justification The benefits of such design of $\mathcal{M}_{\text{rank}}$ are two-folds. (a) The embeddings are created by explicitly taking correlated values into account, so that we can derive comprehensive ranking with contextual information, e.g., we can use the number of kids as additional hints to determine the up-to-date marital status. (b) The encoded vectors $\phi_{t[A]}$ are able to preserve the timeliness (via their distances to the target vector ϕ_A) and meanwhile, reflect the ranking semantics (e.g., the longer timespan from “birth” to “engaged”).

We next briefly elaborate the context-aware embedding and the chronological encoding scheme with the adaptive margin.

Context-aware embedding. To reference correlated values, we treat tuples as sequences and adopt the idea of serialization [43] (so that tuples can be meaningfully ingested by models) to embed values.

Following [43], given a tuple t in D_t , we serialize its values:

$$\text{serialize}(t) = \langle \text{COL} \rangle A_1 \langle \text{VAL} \rangle t[A_1] \dots \langle \text{COL} \rangle A_n \langle \text{VAL} \rangle t[A_n],$$

where $\langle \text{COL} \rangle$ and $\langle \text{VAL} \rangle$ are special tokens, denoting the start of attribute and value, respectively (see Figure 4). This serialization is fed as input to a pre-trained language model $\text{emb}(\cdot)$ to compute a d -dimensional embedding for each A -attribute value, denoted by $\text{emb}(t[A]) \in \mathbb{R}^d$. Besides, we average out the embedding vectors for all $t[A]$ to get a context representation of tuple t , i.e., $\text{emb}(t) = \frac{1}{n} \sum_{i=1}^n \text{emb}(t[A_i])$. Finally for each A -attribute value of t , we get the context-aware embedding of $t[A]$, denoted by $E_{t[A]} \in \mathbb{R}^{2d}$:

$$E_{t[A]} = [\text{emb}(t[A]); \text{emb}(t)],$$

where $[\cdot]$ denotes vector concatenation. In this way, $E_{t[A]}$ embeds

not only the A -attribute value, but also the contextual information from other attribute values, to allow comprehensive ranking. Similarly, a schema embedding for each attribute A is computed: $E_A = \text{emb}(A)$, by feeding the attribute name, e.g., status, to the model.

Chronological encoding with adaptive margins. While pre-trained embeddings are widely adopted to capture semantics, they are not trained to organize temporal orders. Thus we propose chronological encoding to re-organize the embeddings to preserve timeliness. The idea is to use schema embedding as the target and make the embedding of a more current value closer to the target; moreover, instead of ranking in fixed margins, embeddings are ordered adaptively.

Specifically, given the embedding of the A -attribute of t , i.e., $E_t[A]$, we encode it using a context encoder $\text{ENC}_{\text{ctx}}(\cdot)$ as follows:

$$\phi_t[A] = \text{ENC}_{\text{ctx}}(E_t[A]) = \sigma(W_2 * \sigma(W_1 * E_t[A])),$$

where W_1 and W_2 are learnable parameters of the encoder, and σ is the non-linear sigmoid activation function given by $\sigma(x) = \frac{1}{1+e^{-x}}$.

Similarly, the target vector for attribute A is encoded as $\phi_A = \text{ENC}_{\text{attr}}(E_A)$, where $\text{ENC}_{\text{attr}}(\cdot)$ denotes the schema encoder.

To train the encoders with ordered embeddings and adaptive margins, we adopt an attribute-centric adaptive margin-based loss. Given \leq_A in D_t , the loss on A is formulated as follows:

$$\text{loss}(A) = \sum_{(t_1, t_2) \in \leq_A} \left\{ -\tanh(\langle \phi_{t_2}[A], \phi_A \rangle) + \text{ReLU}(\gamma_{t_1, t_2} + \tanh(\langle \phi_{t_1}[A], \phi_{t_2}[A] \rangle)) \right\},$$

where $\langle \cdot, \cdot \rangle$ is the inner product and γ_{t_1, t_2} is the adaptive margin between the two tuples; we set γ_{t_1, t_2} to be $1 - \cos(v_{t_1}[A], v_{t_2}[A])$ in practice, where $v_{t_1}[A]$ and $v_{t_2}[A]$ are the Word2Vec [12] embeddings which characterize the co-occurrence of $t_1[A]$ and $t_2[A]$.

Intuitively, by minimizing the loss, for each training instance $t_1 \leq_A t_2$, (a) we make the encoded vector of the more current value $t_2[A]$ closer to target ϕ_A (the first term) and (b) we ensure an adaptive margin γ_{t_1, t_2} between the two tuples (the second term). In other words, the attribute values in the encoded space are not only arranged chronologically by their distances to ϕ_A , from which temporal orders can be easily derived, their margins are also adaptively determined, to reflect the semantic of timeliness ranking.

Note that w.l.o.g. we follow the common practice in NLP that the schema embeddings will not change after pre-training [3]. To cope with changes, one can use incremental training (see below).

Model training and instance extension. In each round, the creator receives augmented training data $D_{\text{aug}} = (D, \leq'_{A_1}, \dots, \leq'_{A_n}, T, f_{\text{valid}})$, based on which it (a) incrementally trains $\mathcal{M}_{\text{rank}}$ via back-propagation and (b) extends the temporal instance D_t with D_{aug} .

In the first round, D_{aug} is initialized to be the temporal orders constructed by applying those CCs in Σ without timeliness comparison in their preconditions [30]. In the following rounds, D_{aug} is constructed by the critic, based on the result of chasing with CCs.

Specifically, depending on flag f_{valid} in D_{aug} , we have two cases:

(1) If f_{valid} is true, the result of chasing is valid and $\mathcal{M}_{\text{rank}}$ is incrementally trained on the orders in D_{aug} (see below). Moreover, the temporal instance $D_t = (D, \leq_{A_1}, \dots, \leq_{A_n}, T)$ is extended with the ranked pairs in D_{aug} , i.e., for each (t_1, t_2) in \leq'_A of D_{aug} , (t_1, t_2) is added to \leq_A . In this case, we say that (t_1, t_2) becomes *stable*. Once a temporal order becomes stable, it will not be removed from D_t .

(2) If f_{valid} is false, the result of chasing is invalid, i.e., there is an attribute A such that conflicting orders (t_1, t_2) and (t_2, t_1) are both in \leq'_A of D_{aug} i.e., $\{(t_1, t_2), (t_2, t_1)\} \subseteq \leq'_A$, and either $t_1 \prec_A t_2$ or $t_2 \prec_A t_1$. In this case, we decide that either (t_1, t_2) or (t_2, t_1) is added to \leq_A using $\mathcal{M}_{\text{rank}}$, with a higher confidence (and possibly user inspection). Assume w.l.o.g. that (t_1, t_2) is added to \leq_A (i.e., it becomes stable). Then, the creator fine-tunes $\mathcal{M}_{\text{rank}}$ so that t_1 and t_2 are better separated in the encoded space (see below).

Incremental training. The incremental training of $\mathcal{M}_{\text{rank}}$ might lead to the catastrophic forgetting issue [36], i.e., the model might forget some temporal orders learned in prior rounds. To overcome this, we adopt a simple strategy to retain prior knowledge: In each round, $\mathcal{M}_{\text{rank}}$ first makes inference to previously learned orders, and then extracts those that make wrong predictions. Then, we add these orders to the augmented data D_{aug} and then train the model on D_{aug} . In this way, the model is able to learn from previous training instances and alleviate the impact of the catastrophic forgetting issue.

Fine-tuning based on conflicts. Assume w.l.o.g. that (t_1, t_2) is added to \leq_A , since it has a much higher confidence score $\text{conf}(t_1 \leq_A t_2)$ than its conflicting order (t_2, t_1) . We adopt a regularization term in the loss as follows, such that t_1 and t_2 can be better separated:

$$\text{loss}_{\text{reg}}(A) = \text{loss}(A) - \lambda \text{conf}(t_1 \leq_A t_2),$$

where λ is a hyper-parameter, and the second term is the regularization term, penalizing the conflicting order to enlarge their margin. During the incremental training on augmented training data D_{aug} with $f_{\text{valid}} = \text{false}$, we adopt $\text{loss}_{\text{reg}}(A)$ instead of $\text{loss}(A)$.

Monotonicity. One can verify that the number of stable temporal orders in D_t is (strictly) monotonically increasing when more rounds GATE are performed (see a detailed lemma in [1]). The termination of GATE (Theorem 1) partly depends on this monotonicity.

Confidence. Given \leq_A and a tuple pair (t_1, t_2) , $\mathcal{M}_{\text{rank}}$ predicts $(t_1, t_2) \in \leq_A$ if $\tanh(\langle \phi_{t_2}[A], \phi_A \rangle) > \tanh(\langle \phi_{t_1}[A], \phi_A \rangle)$; its confidence indicates how likely $t_1 \leq_A t_2$ holds and it is computed as:

$$\text{conf}(t_1 \leq_A t_2) = \sigma(\tanh(\langle \phi_{t_2}[A], \phi_A \rangle) - \tanh(\langle \phi_{t_1}[A], \phi_A \rangle)),$$

Intuitively, we use $\tanh(\langle \phi_t[A], \phi_A \rangle)$ to measure the “distance” between $\phi_t[A]$ and ϕ_A , where the closer one is more current. The distance gap between $\phi_{t_1}[A]$ and $\phi_{t_2}[A]$ quantifies the confidence: the larger the gap, the larger the confidence, which ranges from 0 to 1.

Example 3: Consider the example in Figure 4, where we focus on attribute status. After creating the context-aware embedding based on a pre-trained model, it chronologically encodes a target vector ϕ_{status} and value vectors for all values, so that they are arranged by their distances to ϕ_{status} . To illustrate, we also label the *unknown* timestamp of each value vector in the figure (e.g., $T_e(\text{Married}) = 2018$). Since ϕ_{Married} is the closest to ϕ_{status} , it is predicted to be the latest value for status and new ranked pairs are constructed accordingly for \leq_{status} , as augmented training data in next round. \square

Remark. Our creator learns temporal orders by utilizing context-aware embedding, chronological encoding and attribute-centric adaptive ranking. However, it does not explicitly take into account of some temporal properties, such as the transitivity. This motivates us to use critic to deduce and justify the temporal orders based on the semantics of the data, as will be presented in the next section.

5 CRITIC

In this section, we develop the critic under GATE for deducing and justifying temporal orders. Taking a temporal instance D_t , the temporal orders $(\leq_{A_1}^M, \dots, \leq_{A_n}^M)$ predicted by the creator and a set Σ of **mined CCs** as input, the critic (a) deduces more ranked pairs $(\leq_{A_1}^\Sigma, \dots, \leq_{A_n}^\Sigma)$ by applying CCs via the *chase*, and (b) constructs the augmented training data D_{aug} and feeds D_{aug} back to the creator.

5.1 Chasing with CCs

We extend the classic chase in [29] for CCs under GATE. In the following, we first specify fixes and ground truth. We then present the chase for CCs, with the Church-Rosser property.

Fixes. We extend temporal orders in D_t by applying CCs in Σ to deduce *fixes*, which are modeled as sets of ranked pairs, denoted by $\bar{U} = (\bar{U}_{A_1}, \dots, \bar{U}_{A_n})$, where each ranked pair (t_1, t_2) in \bar{U}_A is referred to as a *fix* and it means that $t_1 \leq_A t_2$ **or** $t_1 <_A t_2$ is deduced. We only apply a CC if its precondition is satisfied by a collection Γ of “ground truth”. Intuitively, the fixes are logical consequences of Σ and Γ , i.e., as long as the CCs in Σ and Γ are correct, so are the fixes.

Validity. We say \bar{U} is *valid* if it has no conflicting fixes, i.e., there exist no attribute A and tuples t_1, t_2 such that $(t_1, t_2) \in \bar{U}_A$ and $(t_2, t_1) \in \bar{U}_A$ at the same time, **and either** $t_1 <_A t_2$ **or** $t_2 <_A t_1$.

Ground truth. To justify the correctness of fixes, we maintain and employ a collection $\Gamma = (\Gamma_{A_1}, \dots, \Gamma_{A_n})$ of validated data, *enclosed* in \bar{U} . In our setting, block Γ is initialized **by applying CCs in Σ whose preconditions do not involve timeliness comparison (e.g., initial temporal orders with partial timestamps, i.e., φ_3)**, and is iteratively expanded with the temporal orders learned by the creator with confidence above threshold δ or deduced by the chase in the critic.

The chase. Given a temporal instance D_t , the chase deduces fixes by chasing D_t with CCs in Σ and ground truth in Γ . It uses sets $\leq^\Sigma = (\leq_{A_1}^\Sigma, \dots, \leq_{A_n}^\Sigma)$ to keep track of the affected fixes in \bar{U} . Specifically, the i -th chase step of D_t by Σ at $(\bar{U}_i, \leq_i^\Sigma)$ is:

$$(\bar{U}_i, \leq_i^\Sigma) \Rightarrow_{(\varphi, h)} (\bar{U}_{i+1}, \leq_{i+1}^\Sigma),$$

where $\varphi : X \rightarrow p_0$ is a CC in Σ , h is a valuation of φ in \mathcal{D}_t , and the application of (φ, h) should satisfy the following conditions:

- (1) All predicates $p \in X$ are *validated*, i.e., if p is $t[A] \oplus c$ or $t_1[A] \oplus t_2[A]$, then $h \models p$; and if p is $t_1 \leq_A t_2$, then (t_1, t_2) is in \bar{U}_A .
- (2) The consequence $p_0 : t_1 \leq_A t_2$ extends \bar{U}_i to \bar{U}_{i+1} , such that (t_1, t_2) is added to \bar{U}_A of \bar{U}_i ; similarly, p_0 extends \leq_i^Σ to \leq_{i+1}^Σ .

Chasing. Starting from a set \bar{U}_0 of fixes, initialized to be Γ , and an empty \leq_0^Σ , a chasing sequence ξ of D_t by (Σ, Γ) is

$$(\bar{U}_0, \leq_0^\Sigma), \dots, (\bar{U}_k, \leq_k^\Sigma),$$

where $(\bar{U}_i, \leq_i^\Sigma) \Rightarrow_{(\varphi, h)} (\bar{U}_{i+1}, \leq_{i+1}^\Sigma)$ is a valid chase step where a valuation h of φ extends $(\bar{U}_i, \leq_i^\Sigma)$ to $(\bar{U}_{i+1}, \leq_{i+1}^\Sigma)$, i.e., \bar{U}_{i+1} is valid.

The chasing sequence is *terminal* if there **no CC φ in Σ and valuation h of φ** such that (φ, h) could lead to another valid chase step.

A chase sequence ξ terminates in one of the following cases:

- (1) No more CCs in Σ can be applied. If so, we say that ξ is valid, with $(\bar{U}_k, \leq_k^\Sigma)$ as its result.

- (2) Either \bar{U}_0 is invalid or there exist $\varphi, h, \bar{U}_{k+1}$ and \leq_{k+1}^Σ such that $(\bar{U}_k, \leq_k^\Sigma) \Rightarrow_{(\varphi, h)} (\bar{U}_{k+1}, \leq_{k+1}^\Sigma)$ but \bar{U}_{k+1} is invalid. Such ξ is invalid, and the result of the chase is \perp (undefined).

Intuitively, the chase helps us deduce more ranked pairs when it terminates with enriched $(\bar{U}_k, \leq_k^\Sigma)$; moreover, it justifies and explains the learned order if no invalid chase step is taken. When its result is \perp , it detects invalid ranked pairs of the learner.

Example 4: Consider D_t in Figure 1. Assume that $\Sigma = \{\varphi_1, \varphi_6\}$ and $\leq^\Sigma = (\leq_{A_1}^\Sigma, \dots, \leq_{A_n}^\Sigma)$, where each $\leq_{A_i}^\Sigma$ is empty. We initialize \bar{U}_0 and Γ by applying CCs without timeliness comparison, as we did in Example 2, e.g., since $t_1[\text{status}]$ (resp. $t_2[\text{status}]$) is “single” (resp. “married”) in Figure 1, $t_1 \leq_{\text{status}} t_2$ is initialized in Γ by applying φ_2 .

We have the following chase steps of D_t by (Σ, Γ) :

- (1) By applying (φ_4, h_4) , where φ_4 is $t_1 \leq_{\text{status}} t_2 \rightarrow t_1 \leq_{\text{address}} t_2$ and h_4 maps the variables of φ_4 to tuples t_1 and t_2 in D_t , we deduce $t_1 \leq_{\text{address}} t_2$ by the chase step $(\bar{U}_0, \leq_0^\Sigma) \Rightarrow_{(\varphi_4, h_4)} (\bar{U}_1, \leq_1^\Sigma)$, i.e., \bar{U}_1 extends \bar{U}_0 by adding (t_1, t_2) to \bar{U}_{address} ; similar to \leq_1^Σ .
- (2) The chase proceeds to deduce $t_1 \leq_{\text{status}} t_4$ by applying φ_5 .

This chasing sequence is valid since each chase step in the sequence is valid and no more CCs in Σ can be applied anymore. \square

Church-Rosser property. Following [2], we say that chasing with CCs is *Church-Rosser* if for any temporal instance D_t , any set Σ of CCs, any collection Γ of ground truth, all chasing sequences of D_t by (Σ, Γ) are terminal, and all terminal sequences converge at the same result. Below we show that chasing temporal instances with CCs is well-defined and converges at a unique result, no matter what CCs in Σ are used and in what order the CCs are applied.

Corollary 2: Chasing with CCs is Church-Rosser. \square

Proof sketch. CCs can be considered as a special case of entity enhancing rules (REEs) [29], extended with temporal orders \leq_A . We prove Church-Rosser for CCs using a similar argument for REEs in [29], by showing (a) the length of any chasing sequence is bounded and (b) all chasing sequences converge at the same results. \square

5.2 Deduction with the Chase

No matter how desirable, the chase could be expensive if we enumerate valuations of CCs **in an exhaustive manner**. Below we provide an efficient algorithm to implement the chase.

Challenges. A brute-force implementation of the chase is by enumerating the valuation h of each CC φ in Σ . If (φ, h) can be applied, a chase step is performed, until the chase sequence terminates. This method is, however, costly since valuation enumeration is inherently exponential. To tackle this challenge, we develop an efficient algorithm **to implement** the chase; the key idea is to only evoke valuations *pertaining* to the affected fixes in the chase *lazily* (see below).

We assume w.l.o.g. that for each $\varphi : X \rightarrow p_0$ in Σ , X has a predicate p in the form of $t_1 \leq_A t_2$ (e.g., φ_4). For those CCs that do not compare timeliness in their precondition (e.g., φ_1), we apply them in a pre-processing step, to generate initial temporal orders in D_t .

Lazy evocation. To allow lazy evocation, the valuations of CCs in Σ are generated only when they are evoked by some newly deduced orders, instead of constructing all at the beginning of the chase.

Specifically, when a new temporal order o is deduced, we check

each $\varphi : X \rightarrow p_0$ in Σ and evoke a new valuation h of φ if (a) o corresponds to a predicate in X (i.e., o is validated in h); in this case, we say that h is a valuation *pertaining to* o since h is “activated” by o , (b) h has not been evoked before and (c) the order that h deduces, i.e., $h(p_0)$, is not deduced by other valuations before. We maintain designated data structures for checking conditions (a), (b) and (c) efficiently. See [1] for a detailed description of the data structures.

Algorithm. Putting these together, we present Chase in Figure 5. It returns new orders \leq^Σ if the chase is valid, and \perp otherwise.

Chase starts with the initialization (Line 1). (a) Ground truth Γ is initialized with all stable ranked pair in D_t via procedure Initialize (omitted). (b) It initializes \bar{U} and \leq^Σ as stated in Section 5.1 to be Γ and \emptyset , respectively. (c) The set Δ of newly validated orders is initialized to be the newly stable orders in Γ via procedure NewStable (omitted); and they are the temporal orders “triggering” the chase. (d) The set \mathcal{H} of evoked valuations is initialized to be empty.

Then for each order o in Δ , Chase does the following (Line 5-15): (a) Evoke the valuations pertaining to o via the lazy evocation strategy stated above, by calling CCEvoke (omitted), and add them to \mathcal{H} (Line 5). (b) For each valuation h pertaining to o that deduces unknown ranked pair (checked in Line 7-8), mark o as validated in h (Line 9). If all predicates in the precondition of h are validated (Line 10), (h, φ) can be applied and the consequence $t_1 \leq_A t_2$ of h is deduced (Line 11). The ranked pair (t_1, t_2) is added to \bar{U}_A and \leq_A^Σ (Line 12). (c) Check conflicts (Line 13-14): if (t_1, t_2) conflicts with (t_2, t_1) that is already in \bar{U}_A or \leq_A^M , the chase terminates with $\leq^\Sigma = \perp$. In this case, the valuations in \mathcal{H} are kept temporally so that they can be re-used in the next round of GATE when the conflicts are resolved (not shown). (d) Add the newly deduced order to the set Δ^{new} (Line 15) for iteratively processing, by assigning Δ^{new} to Δ (Line 16).

Finally, the result of chasing, \leq^Σ , is returned (Line 17).

Example 5: Recall that φ_4 is $t_1 \leq_{\text{status}} t_2 \rightarrow t_1 \leq_{\text{address}} t_2$ and φ_5 is $t_1 \leq_{\text{status}} t_2 \wedge t_2 \leq_{\text{status}} t_3 \rightarrow t_1 \leq_{\text{status}} t_3$. Let $\Sigma = \{\varphi_4, \varphi_5\}$ and $\Delta = \{t_1 \leq_{\text{status}} t_2\}$. We process $t_1 \leq_{\text{status}} t_2$ in Δ as follows. It first evokes valuations h_4 and h_5 which map the variables of φ_4 and φ_5 to concrete tuples t_1 and t_2 in D_t , with $t_1 \leq_{\text{status}} t_2$ validated. Since the predicate $t_2 \leq_{\text{status}} t_3$ in h_5 is not validated, h_5 is kept in \mathcal{H} for later processing. In contrast, all predicates in the precondition of h_4 are validated and φ_4 deduces $t_1 \leq_{\text{address}} t_2$. Suppose that there is no conflicting order in \bar{U}_{address} and \leq_{address}^M . Then $t_1 \leq_{\text{address}} t_2$ forms a new set Δ and the entire process continues, until Δ is empty. \square

Complexity. The loop of Chase executes at most $\mathcal{O}(|R||D|^2)$ times, since there are at most $|R||D|^2$ temporal orders to be deduced (i.e., the length of any chasing sequence is $\mathcal{O}(|R||D|^2)$). For each temporal order o deduced, we evoke CCs based on o and update the data structures in $\mathcal{O}(c_{\text{val}}|\Sigma|)$ time, where c_{val} denotes the unit cost of constructing the valuations for fixed o and φ , and there are at most $|\Sigma|$ many CCs. Thus, Chase takes at most $\mathcal{O}(c_{\text{val}}|\Sigma||R||D|^2)$ time.

Augmented training data construction. Recall that the result of chasing, denoted by \leq^Σ , is valid or invalid. Based on \leq^Σ , we construct the augmented training data D_{aug} as follows.

(1) If \leq^Σ is valid, both the temporal orders deduced by the chase and predicted by the creator are used to create $D_{\text{aug}} = (D, \leq_{A_1}', \dots, \leq_{A_n}', T, f_{\text{valid}} = \text{true})$ where $\leq_{A_i}' = \leq_{A_i}^M \cup \leq_{A_i}^\Sigma$ ($i \in [1, n]$).

Input: A temporal instance D_t , the set Σ of CCs, the predicted $(\leq_{A_1}^M, \dots, \leq_{A_n}^M)$
Output: The result of chasing, \leq^Σ .

```

1.  $\Gamma := \text{Initialize}(D_t)$ ;  $\bar{U} := \Gamma$ ;  $\leq^\Sigma := \emptyset$ ;  $\Delta = \text{NewStable}(\Gamma)$ ;  $\mathcal{H} := \emptyset$ ;
2. while  $\Delta$  is not empty do
3.    $\Delta^{\text{new}} = \emptyset$ ;
4.   for each  $o \in \Delta$  do
5.      $\mathcal{H} := \mathcal{H} \cup \text{CCEvoke}(D_t, \Sigma, o)$ ;
6.     for each  $h$  of  $\varphi : X \rightarrow t_1 \leq_A t_2$  s.t.  $h$  pertains to  $o$  do
7.       if the order between  $t_1[A]$  and  $t_2[A]$  is already settled then
8.          $\mathcal{H} := \mathcal{H} \setminus \{h\}$ ; continue ;
9.       Mark  $o$  as validated in  $h$ ;
10.      if all predicates in the precondition of  $h$  are validated then
11.         $\mathcal{H} := \mathcal{H} \setminus \{h\}$ ; /*  $t_1 \leq_A t_2$  is a newly deduced order */
12.         $\bar{U}_A := \bar{U}_A \cup (t_1, t_2)$ ;  $\leq_A^\Sigma := \leq_A^\Sigma \cup (t_1, t_2)$ ;
13.        if  $(t_2, t_1) \in \bar{U}_A$  or  $(t_2, t_1) \in \leq_A^M$  then /* conflict */
14.           $\leq^\Sigma := \perp$ ; return  $\leq^\Sigma$ ;
15.         $\Delta^{\text{new}} := \Delta^{\text{new}} \cup \{t_1 \leq_A t_2\}$ ;
16.       $\Delta := \Delta^{\text{new}}$ ;
17. return  $\leq^\Sigma$ ;
```

Figure 5: Procedure Chase

(2) If \leq^Σ is \perp , then there exists conflicting ranked pairs, i.e., both (t_1, t_2) and (t_2, t_1) are in \bar{U}_A or \leq_A^M with $t_1 <_A t_2$ or $t_2 <_A t_1$. In this case, we construct D_{aug} to be $(D, \leq_{A_1}', \dots, \leq_{A_n}', T, f_{\text{valid}} = \text{false})$ where $\leq_{A_i}' = \{(t_1, t_2), (t_2, t_1)\}$ if $A_i = A$ and $\leq_{A_i}' = \emptyset$ otherwise.

As shown in Section 4, the creator fine-tunes its model by using the deduced ranked pairs or the detected conflicts in D_{aug} .

6 EXPERIMENTAL STUDY

Using real-life and synthetic data, we evaluated (1) the effectiveness and (2) the efficiency of GATE for determining temporal orders. We also (3) conducted a case study to showcase the usefulness of GATE.

Experimental settings. We start with the experimental setting.

Datasets. We used three real-life datasets and one synthetic dataset.

(1) Career [37], a benchmark about the careers of football players from FIFA-15 to FIFA-22; it contains 108.5K tuples from 27.2K entities with 20 attributes. We determine the timeliness of potential, position, international reputation and league name. (2) NBA [17], a dataset that encompasses the careers of basketball players; it contains 10.6K tuples with 12 attributes. We determine the currency of team, pts (points) and weight. (3) COM [31], an open-source dataset about self-employed entrepreneurs in Shenzhen. The original dataset contains 1,997,271 tuples with 35 attributes. We removed duplicates and used 8 attributes by excluding date and digital sequences (e.g., the organization code); we derive the timeliness of entrepreneur names from the remaining 1,983,698 tuples. (4) Person, a synthetic dataset of person with 12.3K tuples from 1K entities. Just like Figure 1, we adopted 7 attributes and determine the currency of Address, Status, Kids. Here Person is generated by enforcing CCs (e.g., $\varphi_1\text{-}\varphi_6$) so that the tuples simulate real-world scenarios.

All datasets have ground truth, i.e., all tuples carry timestamps and are grouped by entities. The timestamps of COM are in seconds, while the others are in years; it is reasonable since, e.g., it is uncommon for NBA players to frequently change teams in one year. Since our goal is to deduce the relative order of values, the granularity of timestamps is often not critical. We randomly selected 5% timestamps as the initial partial timestamps and masked the remaining.

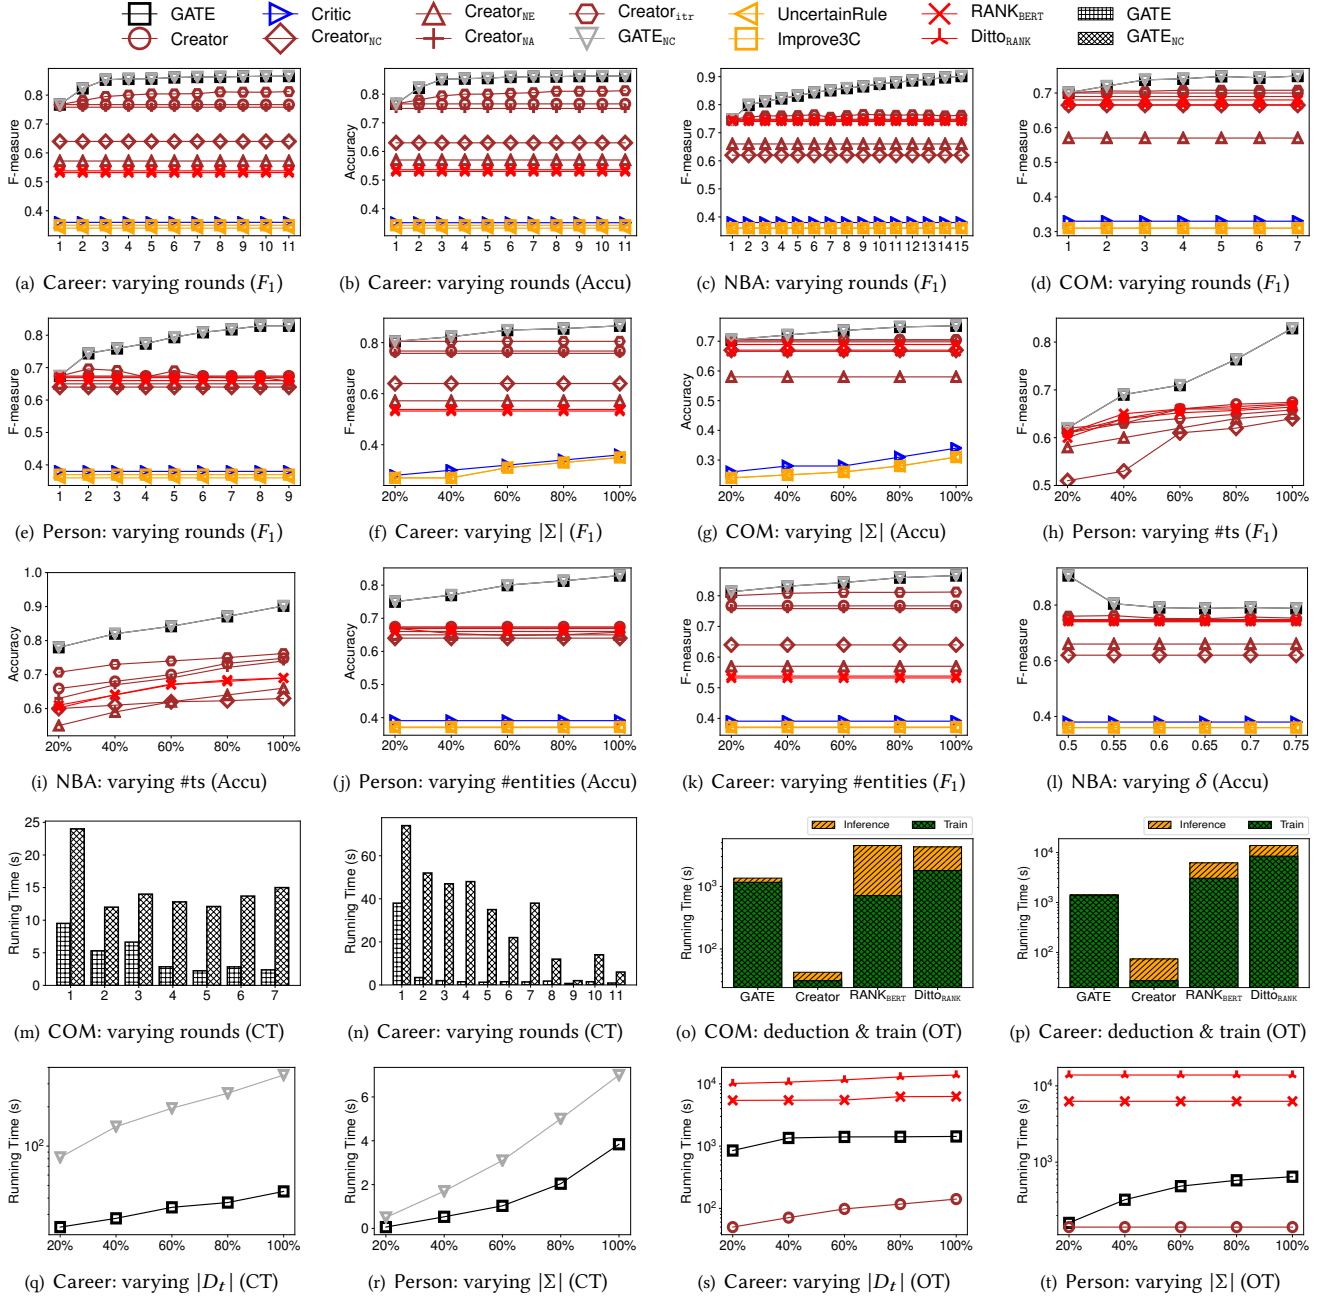


Figure 6: Performance evaluation

Currency constraints. We extended DCFinder [53] to discover CCs as discussed in Section 2. Note that CC discovery is conducted once on each dataset offline. Besides, we manually checked and adjusted the CCs discovered to ensure their correctness. We found 42, 32, 40 and 36 CCs for Career, NBA, COM and Person, respectively.

ML Model. To learn the ranking model $\mathcal{M}_{\text{rank}}$, we used Bert [14] (distilbert) with 768 dimension to initialize the embeddings. We adopted 2 hidden layers in our encoders with sizes 200 and 100 respectively. The margin γ was adaptively computed and the model was trained with 30 epochs using Adam optimizer [35]. The learning rate is $1e-4$. We used 5% data with timestamps as training data.

Baselines. GATE was implemented in Python and we compared it

with the following baselines: (1) Creator, a variant of GATE with the creator only, *i.e.*, it predicts temporal orders using $\mathcal{M}_{\text{rank}}$; (2) Critic, a variant of GATE with the critic only, *i.e.*, it deduces temporal orders by chasing with CCs; (3) Creator_{itr}, a variant of Creator that iteratively update its training data with predicted but unjustified temporal orders. (4) Creator_{NC}, Creator_{NE}, Creator_{NA}, another three variants of Creator that implement $\mathcal{M}_{\text{rank}}$ without contextual information, without chronological encoding, and using regular cross entropy loss instead of adaptive margin-based loss, respectively; (5) GATE_{NC}, a variant of GATE that adopts the brute-force method for the chase, by enumerating all valuations exhaustively.

We also tested (6) UncertainRule [40], which uses uncertain

currency rules to evaluate data currency; (7) Improve3C [17], a data quality framework that combines completeness, consistency and currency [25]; we only compare its **accuracy for currency**; (8) RANK_{Bert} [51], a state-of-the-art ML ranking model based on Bert; and (9) Ditto_{Rank}, a ranking model that first adopts ditto [43] to conduct binary classification on attribute values with their contextual information and then sorts them by the outputted confidences.

Among the baselines, (a) Critic, UncertainRule and Improve3C are rule-based methods, where the rules for UncertainRule and Improve3C are **converted** from the same set of CCs discovered by DCfinder, (b) Creator, Creator_{itr}, Creator_{NA}, Creator_{NC}, Creator_{NE}, RANK_{Bert}, and Ditto_{Rank} are ML methods, and (c) GATE_{NC} is a hybrid method, which produces the same results as GATE. Thus, **we compared GATE_{NC} mostly for efficiency**.

We did the experiments on a single machine powered by 256GB RAM and 32 processors with Intel(R) Xeon(R) Gold 5320 CPU @2.20GHz. We ran each experiment 3 times **and report the average**.

Experimental results. We next report our findings.

Exp-1: Effectiveness. We evaluated the accuracy of GATE using the following metrics. (1) $F\text{-measure} = 2 \times \frac{\text{recall} \times \text{precision}}{\text{recall} + \text{precision}}$ [17, 24, 41], where precision is the ratio of true temporal orders determined correctly to all temporal orders predicated true, and recall is the ratio of true temporal orders predicted correctly to all true temporal orders. (2) Accu, the ratio of correct predictions of the relative rank between each value and the latest value, **for each attribute**. Note that $F\text{-measure}$ is evaluated on the entire dataset (excluding the training data), which is fixed no matter how the parameters change. For the ease of presentation, the accuracy of some baselines are not shown in the figures due to their bad performance.

Rounds. We report the accuracy of GATE from the first round till its termination; at the end of the fixed round, we use the current $\mathcal{M}_{\text{rank}}$ in the creator. By alternating the creator and the critic, we compute the accuracy. As shown in Figures 6(a)-6(e), GATE takes 11, 15, 7 and 9 rounds to terminate on Career, NBA, COM and Person, respectively, *i.e.*, GATE converges quickly. Besides, we find the following.

(1) Although the accuracy of GATE might fluctuate (*e.g.*, Figure 6(d)), which is common in ML models [14], it increases with more rounds in most cases, *e.g.*, $F\text{-measure}$ and Accu increase from 0.767 to 0.866 and 0.766 to 0.862 respectively, after 11 rounds on Career. This is because under our framework, the creator iteratively accumulates more training data from the critic so that the model is better trained with more rounds; meanwhile, with better results predicted by the creator, the critic deduces more orders as augmented training data for the creator in subsequent rounds. Besides, Creator_{itr} also suffers from the issue of accuracy fluctuation (*e.g.*, Figure 6(e)) since its model performance is affected the noisy (unjustified) temporal orders accumulated over rounds. The accuracy of other methods does not depend on rounds, as shown in flat lines in the figures.

(2) GATE outperforms Creator and Critic by 11.4% and 47.4% in $F\text{-measure}$ on average, up to 15.5% and 52.2%, improving both. Under our framework, the creator and critic benefit from each other: (a) Creator produces “hidden” temporal orders for Critic to preform deduction and (b) Critic deduces and justifies the orders, which are in turn provided as augmented training data to Creator; on average,

the critic generates 5733 new training data (tuple pairs) per round on COM, improving Accu of GATE from 0.701 to 0.748 after 5 rounds.

(3) Creator is more accurate than all its variants, *e.g.*, the average $F\text{-measure}$ of Creator is 0.722, as opposed to 0.641, 0.613, and 0.714 by Creator_{NC}, Creator_{NE} and Creator_{NA}, respectively, on Career. Intuitively, (a) without utilizing the contextual information, Creator_{NC} cannot reference correlated attributes; (b) Creator_{NE} has low accuracy with existing embedding models, and (c) compared to the regular cross entropy loss used in Creator_{NA}, the adaptive pairwise ranking loss helps since it considers the semantics in ranking. Moreover, GATE is also 10.1% more accurate than Creator_{itr} on average. This verifies the usefulness of justifying the temporal orders learned by Creator, in order to achieve a better overall accuracy.

(4) The accuracy of GATE is higher than UncertainRule, Improve3C, RANK_{Bert} and Ditto_{Rank}, *e.g.*, the average $F\text{-measure}$ (resp. Accu) of GATE is 0.836 (resp. 0.831) as opposed to 0.344, 0.349, 0.659 and 0.651 (resp. 0.340, 0.342, 0.653 and 0.647) for the four, respectively. This shows the benefits of combining deep learning and logic rules: (a) compared with rule-based methods, GATE can learn from unseen data and has better generalizability; and (b) compared with ML-based methods, GATE is able to justify the reliability of deduction and produces more training data for the model.

Varying $|\Sigma|$. Varying $|\Sigma|$ from 20% to 100%, we evaluated the impact of the number of CCs in Figures 6(f)-6(g). The accuracy of GATE, UncertainRule and Improve3C improves when given more rules. For GATE, the two measures change from 0.805 to 0.866 and 0.705 to 0.752 when $|\Sigma|$ is from 20% to 100%, respectively. Indeed, the critic is able to deduce more orders with more CCs for the creator to fine-tune its model, to get a higher accuracy in an earlier stage.

Varying #ts. We varied the portion of data with initial timestamps, *i.e.*, #ts, from 20% to 100% of the initial training data (the 5% of data with timestamps). Initial timestamps help because (a) the creator has more training instances and can have better initial performance; and (b) the critic gets temporal orders as ground truth to perform deduction at the beginning of the chase. As shown in Figures 6(h) and 6(i), the $F\text{-measure}$ of GATE increases given more data with available timestamps, as expected, *e.g.*, $F\text{-measure}$ increases from 0.62 to 0.829 on Person when #ts changes from 20% to 100%.

Varying #entities. As reported in Figures 6(j) and 6(k), we varied the percentage of entities that are used for the chase from 20% to 100%. As expected, GATE improves its accuracy (both $F\text{-measure}$ and Accu) since with more entities, the critic is able to deduce more temporal orders via the chase, and the creator can have more augmented training data to train the model, achieving higher accuracy. For instance, the Accu of GATE is improved by 6.5% on average.

Varying δ . We next tested the impact of confidence threshold δ . As shown in Figure 6(l), (a) although when δ is small, less confident predictions may appear in subsequent deductions and predictions, more temporal orders could be used for training; (b) when δ is too large, few ranked pairs learned are retained, and hence less augmented training data is returned. Since the creator receives less data to fine-tune its model, the accuracy may not be improved and it converges slowly. When $\delta = 0.5$, GATE has the highest accuracy on NBA. Thus, we set $\delta = 0.5$ as its default; similarly for other

datasets, δ is set to be 0.55.

Exp-2: Efficiency. We next tested GATE and GATE_{NC}, RANK_{Bert} and Ditto_{Rank}. Denoted by CT (resp. OT) the chase time (resp. the overall time, including both model training and rule deduction); for GATE, OT is accumulated over all rounds. We did not report the rule-based methods, which are fast, since they do not need to train models; but as shown in Exp-1, they are not accurate.

Chase time (rounds). We first report CT of GATE and GATE_{NC} in the iterative process. As shown in Figures 6(m) and 6(n), GATE is substantially faster than GATE_{NC} for all rounds, e.g., it is 3.99X and 15.25X faster than GATE_{NC} on average, up to 6.30X and 32.43X, on COM and Career, respectively. The speedup of GATE is due to the lazy evocation strategy we adopted, which accelerates the chase by maintaining designated structures. In contrast, GATE_{NC} enumerates valuations and incurs redundant computation. Since the critic processes of GATE and GATE_{NC} deduce the same results in each round, they terminate with the same total number of rounds.

Overall time. We next report OT in Figure 6(o)-6(p). Although GATE has multiple rounds, its overall time is comparable to most ML methods and even smaller than some of them, e.g., GATE is 6.47X faster than Ditto_{Rank} on average. This shows that GATE only requires small extra cost to achieve the desired improvement in accuracy. In particular, the inference time of GATE is much smaller than those of RANK_{Bert} and Ditto_{Rank}, again justifying the effectiveness of GATE.

Varying $|D_t|$. We evaluated CT (resp. OT) of GATE and GATE_{NC} (resp. ML methods) by varying $|D_t|$ from 20% to 100% in Figure 6(q) (resp. 6(s)). With larger $|D_t|$, all methods take longer, as expected. Nonetheless, GATE is faster than GATE_{NC} when $|D_t|$ gets larger, since GATE maintains structures to avoid recomputation and does deduction pertaining to affected orders, e.g., GATE is 4.80X faster than GATE_{NC} when $|D_t|$ is 100%; the result for OT is consistent.

Varying $|\Sigma|$. We varied $|\Sigma|$ from 20% to 100% on Person in Figure 6(r) and 6(t). As shown there, GATE is 3.76X faster than GATE_{NC} on average, up to 8.33X, which again verifies the effectiveness of lazy evocation. OT of GATE increases as $|\Sigma|$ is larger, since more training data (i.e., temporal orders) is deduced by CCs to train GATE.

Exp-3: Case study. We use Career to illustrate why GATE works.

(1) Initial prediction. In the first round, GATE trains the creator on data with initial timestamps. Due to the limited (5%) training data, its F -measure is only 0.65, few ranked pairs are confident enough, and some pairs are mispredicted. One of confident and correct predictions is $t_1 \leq_{\text{height}} t_2$, where t_1 and t_2 denote the same player with $t_1[\text{height}] = 1.86$ and $t_2[\text{height}] = 1.88$, respectively. However, while this player moved from team PARMA to SPAL, i.e., $t_1 \leq_{\text{team}} t_2$, the creator makes a wrong prediction $t_2 \leq_{\text{team}} t_1$.

(2) Critic helps Creator. After the creator stage, the critic uses CCs to correct mispredicted temporal orders. By applying $\varphi_7 : t_a \leq_{\text{height}} t_b \rightarrow t_a \leq_{\text{team}} t_b$ to the known $t_1 \leq_{\text{height}} t_2$, it deduces $t_1 \leq_{\text{team}} t_2$, correcting the mistake of Creator. Intuitively, φ_7 holds since \leq_{height} is monotonic, and \leq_{height} and \leq_{team} correlate for young players.

Moreover, Critic provides augmented training data to Creator. For instance, if $t_0 \leq_{\text{league_name}} t_1$ and $t_1 \leq_{\text{potential}} t_2$ are in the ground truth, the critic could apply CC $\varphi_8 : t_a \leq_{\text{league_name}} t_b \wedge$

$t_b \leq_{\text{potential}} t_c \wedge t_a[\text{height}] \leq t_c[\text{height}] \rightarrow t_a \leq_{\text{position}} t_c$, and deduces a new pair $t_0 \leq_{\text{position}} t_2$ that is unknown before. Here φ_8 is learned from the data; intuitively, if a player moves to a new league (as indicated by monotonic \leq_{height}) and if his potential changes, his position is likely adjusted, e.g., from LM to CAM. After the first round, the critic creates 100,277 new ranked pairs as augmented training data, and the creator improves its model with the new data.

(3) Creator helps Critic. One should not expect Critic to deduce total temporal orders from limited data with initial timestamps (5% in our setting). Nonetheless, with augmented training data provided by Critic, Creator is able to learn more and more ranked pairs with high confidence. On average it ranks 2843 tuple pairs with high confidence in the first 5 rounds. These ranked pairs are in turn provided to Critic, for Critic to deduce more new ranked pairs.

(4) The creator learns better with more data. ML models are inclined to get more accurate when given more training data. Creator continually receives more augmented training data and incrementally trains its model accordingly. As a consequence, its F -measure increases from 0.65 to 0.74 (resp. 0.81) after the first (resp. last) round.

Summary. We find the following. (1) Combining deep learning and logic rules makes a promising approach to determining temporal orders. GATE achieves the best accuracy, e.g., 0.866 in F -measure on Career, as opposed to 0.35 and 0.36 by the rule-based UncertainRule and Improve3C, and 0.54, and 0.53 by the ML-based RANK_{Bert} and Ditto_{Rank}. (2) GATE only takes 7 rounds to terminate on COM, which has 1,983,698 tuples with 8 attributes. (3) On average, GATE is 47.4% and 11.4% more accurate than Critic and Creator, respectively. That is, the creator and critic indeed benefit each other by learning from unseen data and deducing new training data via logic rules, respectively. (4) GATE beats GATE_{NC} in efficiency (with the same accuracy) by 9.62X on average, up to 19.37X, verifying the effectiveness of our lazy evocation strategy. (5) GATE has competitive overall time against ML methods, e.g., 1359s on COM, as opposed to 4283s by the fastest of them. Although rule-based methods are fast (since they do not train models), their accuracy are much lower than GATE. (6) Our ML model in GATE is more accurate than Creator_{NC}, Creator_{NE}, Creator_{NA} and Creator_{itr} by 19.5%, 22.3%, 12.2% and 10.1%, respectively, verifying the need for context-aware embedding, chronological encoding and adaptive margin.

7 CONCLUSION

The novelty of the work consists of the following. (1) We formulate a new problem for determining the timeliness of attribute values. (2) As a solution to the problem, we propose a creator-critic framework by combining deep learning and logic deduction, for the two to enhance each other. (3) We develop a novel ranking model to learn temporal orders on attribute values. (4) We show how to justify the learned orders, deduce more ranked pairs and provide feedback for the learner, by extending the chase using CCs. The experimental study has verified that GATE is promising in practice.

One topic for future work is to study how to catch conflicts and missing data values by making use of the learned temporal orders. Another topic is to extend CCs of [27] by embedding ranking models as predicates, such that on the one hand, we can improve the accuracy of ranking with additional logic conditions, and on the other

hand, we can interpret ranked orders in terms of logic predicates.

REFERENCES

- [1] 2022. Full version. <https://>.
- [2] Serge Abiteboul, Richard Hull, and Victor Vianu. 1995. *Foundations of Databases*. Addison-Wesley.
- [3] Muhammad Asif Ali, Yifang Sun, Xiaoling Zhou, Wei Wang, and Xiang Zhao. 2019. Antonym-synonym classification based on new sub-space embeddings. In *Proceedings of the AAAI Conference on Artificial Intelligence*, Vol. 33. 6204–6211.
- [4] Marcelo Arenas, Leopoldo Bertossi, and Jan Chomicki. 1999. Consistent Query Answers in Inconsistent Databases. In *PODS*.
- [5] Tobias Bleifuß, Sebastian Kruse, and Felix Naumann. 2017. Efficient Denial Constraint Discovery with Hydra. *PVLDB* 11, 3 (2017), 311–323.
- [6] Philip Bramsen, Pawan Deshpande, Yoong Keok Lee, and Regina Barzilay. 2006. Inducing Temporal Graphs. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. ACL.
- [7] Chris Burges, Tal Shaked, Erin Renshaw, Ari Lazier, Matt Deeds, Nicole Hamilton, and Greg Hullender. 2005. Learning to rank using gradient descent. In *international conference on Machine learning*. 89–96.
- [8] Nathanael Chambers and Dan Jurafsky. 2008. Jointly Combining Implicit Constraints Improves Temporal Ordering. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)* (Honolulu, Hawaii). ACL, 698–706.
- [9] Olivier Chapelle and Yi Chang. 2011. Yahoo! learning to rank challenge overview. In *Proceedings of the learning to rank challenge*. PMLR, 1–24.
- [10] Peter Christen and Ross W. Gayler. 2013. Adaptive Temporal Entity Resolution on Dynamic Databases. In *PAKDD*. Springer.
- [11] Xu Chu, Ihab F. Ilyas, and Paolo Papotti. 2013. Discovering Denial Constraints. *PVLDB* 6, 13 (2013), 1498–1509.
- [12] Kenneth Ward Church. 2017. Word2Vec. *Natural Language Engineering* 23, 1 (2017), 155–162.
- [13] E. F. Codd. 1979. Extending the Database Relational Model to Capture More Meaning. *TODS* 4, 4 (1979), 397–434.
- [14] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. In *NAACL-HLT*. 4171–4186.
- [15] Ioannis Dikeoulas, Saadullah Amin, and Günter Neumann. 2022. Temporal Knowledge Graph Reasoning with Low-rank and Model-agnostic Representations. *CoRR* abs/2204.04783 (2022).
- [16] Xiaou Ding, Hongzhi Wang, Yitong Gao, Jianzhong Li, and Hong Gao. 2017. Efficient currency determination algorithms for dynamic data. *Tsinghua Science and Technology* 22, 3 (2017), 227–242.
- [17] Xiaou Ding, Hongzhi Wang, Jiaxuan Su, Jianzhong Li, and Hong Gao. 2018. Improve3c: Data cleaning on consistency and completeness with currency. *arXiv preprint arXiv:1808.00024* (2018).
- [18] Xiaou Ding, Hongzhi Wang, Jiaxuan Su, Muxian Wang, Jianzhong Li, and Hong Gao. 2020. Leveraging Currency for Repairing Inconsistent and Incomplete Data. *TKDE* (2020).
- [19] Aswathy Divakaran and Anuraj Mohan. 2020. Temporal Link Prediction: A Survey. *New Gener. Comput.* 38, 1 (2020), 213–258. <https://doi.org/10.1007/s00354-019-00065-z>
- [20] Xuliang Duan, Bing Guo, Yan Shen, Yuncheng Shen, Xiangqian Dong, and Hong Zhang. 2020. Research on Parallel Data Currency Rule Algorithms. In *International Conference on Information Science and System*. 24–28.
- [21] Kevin K Duh. 2009. *Learning to rank with partially-labeled data*. University of Washington.
- [22] Exasol. 2020. Exasol Research Finds 58% of Organizations Make Decisions Based on Outdated Data. <https://www.exasol.com/news-exasol-research-finds-organizations-make-decisions-based-on-outdated-data/>.
- [23] Wenfei Fan, Floris Geerts, Xibei Jia, and Anastasios Kementsietsidis. 2008. Conditional functional dependencies for capturing data inconsistencies. *TODS* 33, 2 (2008), 6:1–6:48.
- [24] Wenfei Fan, Floris Geerts, Nan Tang, and Wenyuan Yu. 2013. Inferring data currency and consistency for conflict resolution. In *ICDE*. IEEE, 470–481.
- [25] Wenfei Fan, Floris Geerts, Nan Tang, and Wenyuan Yu. 2014. Conflict resolution with data currency and consistency. *Journal Data and Information Quality (JDIQ)* 5, 1-2 (2014), 6:1–6:37.
- [26] Wenfei Fan, Floris Geerts, and Jef Wijsen. 2011. Determining the currency of data. In *PODS*. ACM.
- [27] Wenfei Fan, Floris Geerts, and Jef Wijsen. 2012. Determining the Currency of Data. *TODS* 37, 4 (2012), 25:1–25:46.
- [28] Wenfei Fan, Ruochun Jin, Ping Lu, Chao Tian, and Ruiqi Xu. 2022. Towards Event Prediction in Temporal Graphs. *PVLDB* 15, 9 (2022), 1861–1874.
- [29] Wenfei Fan, Ping Lu, and Chao Tian. 2020. Unifying Logic Rules and Machine Learning for Entity Enhancing. *Sci. China Inf. Sci.* 63, 7 (2020).
- [30] Wenfei Fan, Chao Tian, Yanghao Wang, and Qiang Yin. 2021. Parallel Discrepancy Detection and Incremental Detection. *PVLDB* 14, 8 (2021), 1351–1364.
- [31] Shenzhen Municipal Government. 2022. Self-employed Entrepreneurs. https://opendata.sz.gov.cn/data/dataSet/toDataDetails/29200_01300931.
- [32] Tanya Goyal and Greg Durrett. 2019. Embedding Time Expressions for Deep Temporal Ordering Models. In *Conference of the Association for Computational Linguistics (ACL)*. ACL.
- [33] Shuguang Han, Xuanhui Wang, Mike Bendersky, and Marc Najork. 2020. Learning-to-Rank with BERT in TF-Ranking. *arXiv preprint arXiv:2004.08476* (2020).
- [34] Yujing Hu, Qing Da, Anxiang Zeng, Yang Yu, and Yinghui Xu. 2018. Reinforcement learning to rank in e-commerce search engine: Formalization, analysis, and application. In *SIGKDD*. 368–377.
- [35] Diederik P. Kingma and Jimmy Ba. 2015. Adam: A Method for Stochastic Optimization. In *ICLR (Poster)*.
- [36] James Kirkpatrick, Razvan Pascanu, Neil C. Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. 2016. Overcoming catastrophic forgetting in neural networks. *CoRR* abs/1612.00796 (2016).
- [37] Stefano Leone. 2022. FIFA 22 complete player dataset. <https://www.kaggle.com/stefanoleone992/fifa-22-complete-player-dataset>.
- [38] Furong Li, Mong-Li Lee, Wynne Hsu, and Wang-Chiew Tan. 2015. Linking Temporal Records for Profiling Entities. In *SIGMOD*. ACM, 593–605.
- [39] Mohan Li and Jianzhong Li. 2016. A minimized-rule based approach for improving data currency. *J. Comb. Optim.* (2016), 812–841.
- [40] Mohan Li, Jianzhong Li, Siyao Cheng, and Yanbin Sun. 2018. Uncertain rule based method for determining data currency. *IEICE TRANSACTIONS on Information and Systems* 101, 10 (2018), 2447–2457.
- [41] Mohan Li and Yanbin Sun. 2018. Currency Preserving Query: Selecting the Newest Values from Multiple Tables. *IEICE TRANSACTIONS on Information and Systems* 101, 12 (2018), 3059–3072.
- [42] Pei Li, Xin Luna Dong, Andrea Maurino, and Divesh Srivastava. 2011. Linking Temporal Records. *PVLDB* 4, 11 (2011), 956–967.
- [43] Yuliang Li, Jinfeng Li, Yoshihiko Suhara, AnHai Doan, and Wang-Chiew Tan. 2020. Deep Entity Matching with Pre-Trained Language Models. *PVLDB* 14, 1 (2020), 50–60.
- [44] Yu Liang, Xuliang Duan, Yuanjun Ding, Xifeng Kou, and Jingcheng Huang. 2019. Data Mining of Students' Course Selection Based on Currency Rules and Decision Tree. In *International Conference on Big Data and Computing*. 247–252.
- [45] Ashley Little. 2020. Outdated Data: Worse Than No Data? <https://info.aldensys.com/joint-use/outdated-data-is-worse-than-no-data#:~:text=Obsolete%20data%20about%20the%20condition,too%20old%20to%20be%20reliable>.
- [46] Tie-Yan Liu. 2009. Learning to Rank for Information Retrieval. *Found. Trends Inf. Retr.* 3, 3 (2009), 225–331. <https://doi.org/10.1561/15000000016>
- [47] Tie-Yan Liu. 2010. Learning to rank for information retrieval. In *SIGIR*.
- [48] Ester Livshits, Alireza Heidari, Ihab F. Ilyas, and Benny Kinfelfeld. 2020. Approximate Denial Constraints. *PVLDB* 13, 10 (2020), 1682–1695.
- [49] Qiang Ning, Zhili Feng, and Dan Roth. 2017. A Structured Learning Approach to Temporal Relation Extraction. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*. ACL, 1027–1037.
- [50] Qiang Ning, Hao Wu, Haoruo Peng, and Dan Roth. 2018. Improving Temporal Relation Extraction with a Globally Acquired Statistical Resource. In *Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT*. ACL, 841–851.
- [51] Rodrigo Frassetto Nogueira and Kyunghyun Cho. 2019. Passage Re-ranking with BERT. *CoRR* abs/1901.04085 (2019).
- [52] Rama Kumar Pasumarthi, Sebastian Bruch, Xuanhui Wang, Cheng Li, Michael Bendersky, Marc Najork, Jan Pfeifer, Nadav Golbandi, Rohan Anil, and Stephan Wolf. 2019. TF-ranking: Scalable tensorflow library for learning-to-rank. In *SIGKDD*. 2970–2978.
- [53] Eduardo H. M. Pena, Eduardo Cunha de Almeida, and Felix Naumann. 2019. Discovery of Approximate (and Exact) Denial Constraints. *PVLDB* 13, 3 (2019), 266–278.
- [54] Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. Deep contextualized word representations. In *NAACL*.
- [55] Royal Mail. 2018. Dynamic Customer Data in a Digital World: Data Services Insight Report. <https://www.royalmail.com/business/system/files/royal-mail-data-services-insight-report-2018.pdf>.
- [56] Ali Sadeghian, Mohammadreza Armandpour, Anthony Colas, and Daisy Zhe Wang. 2021. ChronoR: Rotation Based Temporal Knowledge Graph Embedding. In *AAAI*. AAAI Press, 6471–6479.
- [57] Fereidoon Sadri and Jeffrey D. Ullman. 1980. The Interaction between Functional Dependencies and Template Dependencies. In *SIGMOD*.
- [58] Yi Tay, Minh C Phan, Luu Anh Tuan, and Siu Cheung Hui. 2017. Learning to rank question answer pairs with holographic dual lstm architecture. In *SIGIR*. 695–704.
- [59] Julien Tourille, Olivier Ferret, Aurélie Névoul, and Xavier Tannier. 2017. Neural Architecture for Temporal Relation Extraction: A Bi-LSTM Approach for Detecting Narrative Containers. In *ACL*. ACL, 224–230.
- [60] Rakshit Trivedi, Hanjun Dai, Yichen Wang, and Le Song. 2017. Know-Evolve: Deep Temporal Reasoning for Dynamic Knowledge Graphs. In *International*

- Conference on Machine Learning (ICML)*, Vol. 70. PMLR, 3462–3471.
- [61] Hongzhi Wang, Xiaoou Ding, Jianzhong Li, and Hong Gao. 2018. Rule-based entity resolution on database with hidden temporal information. *TKDE* 30, 11 (2018), 2199–2212.
 - [62] Jun Xu, Xiangnan He, and Hang Li. 2020. Deep Learning for Matching in Search and Recommendation. *Found. Trends Inf. Retr.* 14, 2-3 (2020), 102–288.
 - [63] Jing Yao, Zhicheng Dou, Jun Xu, and Ji-Rong Wen. 2021. RLPS: A Reinforcement Learning-Based Framework for Personalized Search. *TOIS* 39, 3 (2021), 1–29.
 - [64] Jingran Zhang, Fumin Shen, Xing Xu, and Heng Tao Shen. 2020. Temporal Reasoning Graph for Activity Recognition. *IEEE Trans. Image Process.* (2020).
 - [65] Meng Zhang, Yang Liu, Huanbo Luan, and Maosong Sun. 2016. Listwise ranking functions for statistical machine translation. *IEEE/ACM Transactions on Audio, Speech, and Language Processing* 24, 8 (2016), 1464–1472.

Input: A temporal instance D_t , the set Σ of CCs, the set $(\leq_{A_1}^M, \dots, \leq_{A_n}^M)$ predicted by the creator, the global structures $GS = (RHS, \mathcal{H}, \mathcal{I}, M)$.

Output: The result of chasing, \leq^Σ .

1. $\Gamma := \text{Initialize}(D_t); \bar{U} := \Gamma; \leq^\Sigma := \emptyset; \Delta = \text{NewStable}(\Gamma);$
2. **while** Δ is not empty **do**
3. $\Delta^{\text{new}} = \emptyset;$
4. **for each** $o \in \Delta$ **do**
5. $\mathcal{H} := \mathcal{H} \cup \text{CCEvoke}(D_t, \Sigma, o, GS);$ Update \mathcal{I} and M ;
6. **for each** $h \in \mathcal{I}[o]$ where h deduces $t_1 \leq_A t_2$ **do**
7. **if** $(t_1, t_2, A) \notin \text{RHS}$ **then**
8. Remove h from \mathcal{H} and \mathcal{I} ; **continue** ;
9. Mark o as validated in h ;
10. **if** h is complete **then** /* $t_1 \leq_A t_2$ is a newly deduced order */
11. Remove h from \mathcal{H} and \mathcal{I} ;
12. $\bar{U}_A := \bar{U}_A \cup (t_1, t_2); \leq_A^\Sigma := \leq_A^\Sigma \cup (t_1, t_2);$
13. **if** $(t_2, t_1) \in \bar{U}_A$ or $(t_2, t_1) \in \leq_A^M$ **then** /* conflict */
14. $\leq^\Sigma := \perp$; **return** \leq^Σ ;
15. $\Delta^{\text{new}} := \Delta^{\text{new}} \cup \{t_1 \leq_A t_2\};$
16. $\Delta := \Delta^{\text{new}};$
17. **return** \leq^Σ ;

Figure 7: Procedure Chase (with data structures)

APPENDIX A: DISCOVERY OF CURRENCY CONSTRAINTS

As noted in [27], CCs can be considered as a special case of denial constraints (DCs) introduced in [4], extended with temporal orders \leq_A . Several algorithms have been developed for discovering DCs, e.g., FastDC [11], Hydra [5], DCFinder [53] and ADCMiner [48].

We extend DCFinder [53] for discovering CCs as follows. (1) We support timeliness comparisons, i.e., $t_1 \leq_A t_2$, by adding them to the evidence set used in [53] or transforming categorical values to numerical ones using a mapping function $f_{\text{map}}(\cdot)$, such that timeliness is preserved, e.g., given $t_1 \leq_A t_2$, we ensure $f_{\text{map}}(t_1[A]) \leq f_{\text{map}}(t_2[A])$. (2) We restrict our evidence set on tuples with the same EIDs, instead of using a global evidence set. This accelerates CCs discovery, since CCs are only defined on tuple variables that refer to the same entity. (3) We revise DCFinder by always selecting $t_1[\text{EID}] = t_2[\text{EID}]$ as the first predicate. Note that CCs could also be added manually, e.g., to ensure the transitivity of the data.

APPENDIX B: MONOTONICITY

Monotonicity. We next show that the number of stable temporal orders in D_t is (strictly) monotonically increasing when more rounds GATE are performed (Lemma 3). The termination of GATE (Theorem 1) partly depends on this lemma. Lemma 3 directly follows from the way we expand stable temporal orders.

Lemma 3: For temporal instances D_t^{j-1} and D_t^j in the j -th round of GATE before and after the extension, respectively, i.e., $D_t^j = \text{Extend}(D_t^{j-1}, D_{\text{aug}})$, the following holds:

$$(a) \forall i \in [1, n], \leq_{A_i}^{j-1} \subseteq \leq_{A_i}^j \text{ and } (b) \exists i^* \in [1, n], \leq_{A_{i^*}}^{j-1} \subset \leq_{A_{i^*}}^j$$

where $\leq_{A_i}^{j-1}$ and $\leq_{A_i}^j$ are the orders in D_t^{j-1} and D_t^j , respectively. \square

Proof. By the way we expand stable temporal orders, there are two cases: (1) If f_{valid} is true, the result of chasing is valid. Then at least one non-empty \leq_A' of D_{aug} will be used to extend \leq_A , whose size

strictly increases. (2) If f_{valid} is false, the result of chasing is invalid, i.e., there is an attribute A such that $t_1[A] \neq t_2[A]$ but conflicting orders (t_1, t_2) and (t_2, t_1) are both in \leq_A' of D_{aug} . Either (t_1, t_2) or (t_2, t_1) is added to \leq_A , resulting an increased size of \leq_A . Once a temporal order $t_1 \leq_A t_2$ is added to \leq_A , it will not be removed. \square

APPENDIX C: STRUCTURES AND STRATEGIES

To support lazy evocation of valuations, we employ the following:

(1) A set RHS of triples, where each triple (t_1, t_2, A) in RHS indicates that the order between $t_1[A]$ and $t_2[A]$ is not settled, i.e., neither $t_1 \leq_A t_2$ nor $t_2 \leq_A t_1$ is stable in D_t yet. Note that we only apply a CC if its consequence has a corresponding triple in RHS. By ensuring this, the length of a chasing sequence is bounded by $O(|\text{RHS}|)$.

(2) A set \mathcal{H} of *partial valuations*. Specifically, a valuation h of CC $\varphi : X \rightarrow p_0$ is said to be *partial* if some predicates in X are validated, while others are not. If all predicates in X are validated, h becomes *complete* and we can deduce temporal orders by applying (φ, h) . The set \mathcal{H} is maintained to avoid repeated predicate validation.

(3) An index \mathcal{I} for the partial valuations in \mathcal{H} , i.e., for each temporal order o , $\mathcal{I}[o]$ maintains the partial valuations h of $\varphi : X \rightarrow p_0$ in \mathcal{H} such that there exists a predicate p in X and $o = h(p)$. By maintaining \mathcal{I} , every time a temporal order o is deduced, we can efficiently locate the valuations affected by o in $\mathcal{I}[o]$, without scanning the entire Σ . Besides, an inverted index is also built for each h so that once h is removed from \mathcal{H} , \mathcal{I} can be updated efficiently.

(4) A set M of triples, where each triple (t_1, t_2, φ) indicates that the ranked pair (t_1, t_2) has been used to evoke the valuations for φ before and thus those valuations will not be evoked again.

Moreover, we adopt the following strategies.

(5) Lazy evocation, where valuations of CCs in Σ are constructed if they are evoked by some newly deduced orders, instead of being generated all at the beginning of the chase. Specifically, when a new temporal order o is deduced, we check each $\varphi : X \rightarrow p_0$ in Σ and evoke a new partial valuation h of φ if o corresponds to a predicate in X (i.e., o is validated in h) and h has not been evoked before (checked by M). Such h can only be evoked if the temporal order it deduces, i.e., $h(p_0)$, is not deduced by other valuations before (checked by RHS).

Algorithm. Putting these together, we present Chase in Figure 7 (a complete version of Figure 5, with data structures incorporated). It returns new orders \leq^Σ if the chase is valid, and \perp otherwise.

Chase starts with the initialization (Line 1). (a) Ground truth Γ is initialized with all stable ranked pair in D_t via procedure Initialize (omitted). (b) \bar{U} and \leq^Σ are initialized as stated in Section 5.1 to be Γ and \emptyset , respectively. (c) The set Δ of newly validated orders is initialized to be the newly stable orders in Γ via procedure NewStable (omitted), and they are the temporal orders “triggering” the chase.

Then for each order o in Δ , Chase does the following (Line 5-15): (a) Evoke the valuations h based on o via the lazy evocation strategy stated above, by calling CCEvoke (omitted), and add h to \mathcal{H} (Line 5). (b) For each valuation h in $\mathcal{I}[o]$, mark o as validated in h (Line 9). If h becomes complete (Line 10), the consequence $t_1 \leq_A t_2$ of h is deduced, and the ranked pair (t_1, t_2) is added to \bar{U}_A and \leq_A^Σ (Line 12). (c) Check conflicts (Line 13-14): if (t_1, t_2) conflicts with (t_2, t_1) that is already in \bar{U}_A or \leq_A^M , the chase terminates with $\leq^\Sigma = \perp$.

In this case, the partial valuations and the temporal orders that have been examined and deduced are kept temporally in the global structures so that they can be re-used in the next round of GATE when the conflicts are resolved (not shown). (d) Maintain the global structures in the three cases below: (i) if new valuation h is evoked by o (Line 5), \mathcal{I} is updated accordingly and M is also updated such that h will not be evoked again; (ii) if h becomes useless, *i.e.*, the

ranked pair it deduces is no longer in RHS (Line 7-8), h is removed from \mathcal{H} and \mathcal{I} ; and (iii) if h becomes complete, we deduce new orders, namely $t_1 \leq_A t_2$, by applying h ; then h is removed from \mathcal{H} and \mathcal{I} (Line 11). (e) Add the newly deduced order to the set Δ^{new} (Line 15) for iteratively processing, by assigning Δ^{new} to Δ (Line 16).

Finally, the result of chasing, \leq^Σ , is returned (Line 17).